

A Tutorial in Dynamic Systems Simulation and Modelling With R, ctsem, and lme4: Couples' Affect Dynamics Over Time

Charles C. Driver and Michael Aristodemou
Department of Psychology, University of Zurich

Understanding how people change is a fundamental goal of psychological science. However, translating complex ideas about psychological dynamics into formal models can be challenging without tools that make the modelling assumptions explicit. In this tutorial, we introduce an iterative workflow for dynamic systems modelling in R: theory is translated into a generative model, the model is simulated with understandable functions, a corresponding statistical model is estimated, model-implied behavior is checked, and the specification is revised, simplified, or extended. We begin with familiar linear models in lme4 and gradually transition to ctsem, which allows us to estimate models with state-dependent change, random fluctuation distinct from measurement error, covariate effects, interactions, and external inputs. These modelling concepts are illustrated using a running example of affect dynamics in couples therapy, while also noting how the same model components apply to other psychological domains. Our aim is to provide a general framework for understanding dynamic systems modelling and to encourage further exploration of theory-driven statistical approaches in psychological research.

Impact Statement

This tutorial shows how psychological researchers can use R and the ctsem package to build and understand dynamic systems models while keeping the modelling assumptions visible. By working from theory to simulations, estimation, checking, and revision, we aim to facilitate the translation of theoretical concepts into empirical models and improve understanding of psychological processes and their interrelations over time.

Keywords: dynamic systems, continuous time, hierarchical modelling, longitudinal data analysis, affect dynamics, ctsem, lme4, state-space models

Dynamic systems theory provides a framework for understanding how psychological constructs evolve over time (van Geert, 2011). To turn conceptual ideas about dynamic sys-

tems into insight about mechanisms of change, theory must be translated into empirically testable models, which when combined with data can in turn suggest updates to theory (Borsboom et al., 2021). Yet, this translation from conceptual ideas to testable statistical models can be difficult without tools that make the steps transparent.

 Charles C. Driver

 Michael Aristodemou

An earlier version of this manuscript is available as a PsyArXiv preprint at https://osf.io/preprints/psyarxiv/4q9ex_v3. The accompanying OSF project, including manuscript materials and code, is available at <https://osf.io/xh6ue/>. The examples use simulated data generated for the tutorial. The authors declare that they have no conflicts of interest to disclose. Author roles were classified using the Contributor Role Taxonomy (CRediT; <https://credit.niso.org/>) as follows: Charles C. Driver: conceptualization, methodology, software, validation, writing, visualization; Michael Aristodemou: formal analysis, writing, visualization

Correspondence concerning this article should be addressed to Charles C. Driver, Department of Psychology, University of Zurich, Binzmühlestrasse 14, Box 28, Zurich, Zurich 8057, Switzerland, Email: charles.driver@psychologie.uzh.ch

This tutorial is designed to help psychological researchers bridge the gap between theoretical concepts in dynamic systems and their practical implementation using statistical modelling. The central contribution is the components of a workflow: translate theory into a generative model, simulate that model with transparent R functions, estimate a corresponding statistical model, check whether the model-implied behavior matches the intended theory and observed data, and then revise, simplify, or extend the specification. We begin with familiar linear mixed-effects models using the R package lme4 (Bates et al., 2015), then gradually move into the dynamic systems and continuous-time domain using the package ctsem (Driver & Voelke, 2018a). This progression highlights how modelling choices can evolve alongside theory, allowing increasingly realistic and nuanced representations of psycho-

logical change.

A key strength of *ctsem* is that it can accommodate this full modelling continuum. Starting from basic linear growth models, one can move through discrete- and continuous-time versions of classic frameworks such as the Cross-Lagged Panel Model (Kenny, 2005), the random intercept extension of it (Hamaker et al., 2015), and related forms such as network vector autoregression (Epskamp, 2020). From there, one can extend *ctsem* to models with considerably more complexity. These extensions include complex variation across individuals, oscillations, and measurement or dynamics parameters that shift over time or in response to moderators and interventions. In this way, *ctsem* provides a flexible environment for building models up step by step: researchers can begin with simple forms to establish a baseline, then incorporate the dynamics that theory suggests, then perhaps go further with data-driven model and theory development. While this tutorial emphasises the *ctsem* package, and specific syntax and package issues are discussed, the principles demonstrated are general and can be applied to other packages and modelling approaches.

Continuous-time representations are particularly valuable for theory-oriented model development, because of the more natural link between continuous-time parameters and typical theoretical interests (Aalen et al., 2016; Driver & Tomasik, 2023; Ryan & Hamaker, 2022). Because many psychological processes — such as emotion regulation, stress responses, symptom fluctuation, learning, physiological regulation, or cognitive fatigue — are thought to evolve continuously rather than only when measured, continuous-time models provide a natural representation. Parameters such as drift (state dependence), diffusion (system noise), and external inputs can be aligned much more closely with mechanistic interpretations, instead of as artifacts of the measurement schedule. Continuous-time models are also useful when observation intervals vary or when the time scale of change is itself substantively important. This interpretability makes continuous-time models especially valuable for theory-oriented research, as it ensures that any constraints imposed on the model (e.g., fixing the effect of one variable on another to zero) are meaningfully interpretable with respect to theory (Driver, 2025).

Simulation is a core element of this workflow because it makes the consequences of theoretical assumptions visible before empirical estimation. All simulations in this tutorial are pieced together from simpler R functions, to facilitate a clear link between the model and the simulation. It is of course important to recognise that simulated data will not resemble every empirical complication, but to help readers understand how model parameters generate trajectories, when different sources of variability have different implications, and how a fitted model can fail even when it explains a large amount of variance. Simulation helps build the basis of understanding necessary for applying the principles of dy-

namic systems modelling to inevitably more complex empirical data, and can also help researchers determine approximate requirements in terms of data density, timing, and reliability, under specific assumptions (however, this is not a tutorial on power analysis or sample size determination).

Boundary Conditions

Continuous-time dynamic systems models are most useful when the research question concerns change over time, when theory implies that current states influence subsequent change, when observation timing matters, or when parameters can be linked to interpretable features of an evolving system. They can be especially valuable for intensive longitudinal data, irregularly spaced observations, intervention or event processes, dyadic and networked systems, and questions about heterogeneity in dynamic regulation. They are not automatically preferable, however. Simpler descriptive, mixed-effects, or discrete-time models may be better choices when the research question is not dynamic, when the time scale is not theoretically meaningful, when measurements are too sparse or unreliable to identify the proposed dynamics, or when the parameters of a more complex model would not be substantively interpretable.

Model complexity therefore has to be calibrated in both directions. A model can be too complex for the data, leading to weak identification, unstable estimates, or estimation failures. But a model can also be too simple: distinct influences may be bundled into a single parameter and then interpreted as if that parameter represented one psychological process. In practice, decisions should triangulate theory, measurement design, observation density, timing variation, model-implied behavior, diagnostics, and parsimony. Important empirical complications not fully treated here include missing-not-at-random, nonstationarity, floor or ceiling effects, non-Gaussian outcomes, and many other forms of misspecification.

Scope

In this tutorial we first focus on developing intuition for continuous-time dynamic systems modelling by building up complexity with a running example, then later address model evaluation and improvement. The example we use is the impact of couples therapy on affect dynamics over time, which is grounded in empirical research showing that partners' affective states are dynamically interdependent, becoming synchronized through processes of mutual influence (Butler & Randall, 2013). Such coupling can either amplify or dampen emotional experiences, with implications for individual well-being and relationship stability. Using this context, we progressively build from basic assumptions of linear change to models incorporating feedback loops, random fluctuations, dyadic interdependence, and the effects of therapeutic inputs. Through this step-by-step approach, we introduce different

models of affective processes that reflect increasing levels of complexity:

1. **Steady (Linear) Change:** Affect changes at a constant rate from its starting level between observations—the simplest step beyond a flat baseline. The same structure arises in many other settings, such as developmental change, treatment response, learning, or practice effects, whenever a simple model of systematic change over time is needed.
2. **Nonlinear Change (Discrete Time):** The rate of change now depends on affect at the preceding observation, so trajectories can bend rather than grow in a straight line. The same logic applies, for example, to symptom recovery, skill acquisition, or fatigue when becomes easier or more difficult at different levels.
3. **Nonlinear Change with Random Fluctuations (Continuous Time):** Change now unfolds continuously, with the rate depending on current (rather than past) affect. Random fluctuations represent unmodelled influences that perturb the process between observations, separating measurement error from process noise – useful in any context where the observations do not perfectly reflect the underlying process of interest. This section include a conceptual introduction to continuous-time dynamic modelling.
4. **Couple Dynamics with Self and Partner Effects (Continuous Time):** A multivariate extension in which each partner's rate of change depends on their own affect and their partner's. The same coupling or cross-effects structure could describe, among other cases, linked symptoms, cognition and physiology, parent–child dynamics, or interacting behaviors.
5. **Couple Dynamics with Individual Differences and Time-Independent Moderation (Continuous Time):** People differ in their dynamic parameters, and a stable between-person variable (here, time spent together) moderates partner coupling. Such moderation questions arise when stable traits, diagnoses, treatment groups, age, or context relate to differences in dynamic parameters.
6. **Couple Dynamics with Individual Differences and Time-Dependent Moderation (Continuous Time):** Partner coupling depends on whether the couple is communicating at that moment – a time-varying moderator switches a dynamic parameter on or off. Similar logic applies whenever an observed time-varying condition modifies a dynamic parameter; for example, nicotine replacement may dull the link between certain behaviours or contexts and craving.
7. **Nonlinear Change with a Transient External Shock (Continuous Time):** A therapy session delivers an instantaneous input when it occurs; the effect then propagates and dissipates through the dynamics. Transient

inputs can likewise represent, for example, acute stressors, prompts, medication doses, feedback, or experimental manipulations.

8. **Nonlinear Change with a Persistent External Shock (Continuous Time):** Therapy shifts the process parameters themselves, leaving a lasting change in regulation or level. Persistent inputs could also model, for example, learning, cumulative treatment response, habit formation, or enduring environmental change.

These different models provide researchers with a structured approach to modelling psychological dynamic, allowing for the representation of both complex within-person change and between-person heterogeneity. While it is not necessary for the reader to engage thoroughly with every modelling example, such modelling competencies are an important basis for considering the final, and in some sense most important section of the tutorial, addressing model evaluation in the dynamic systems context. In this section we discuss some necessary checks to ensure the fitted model is appropriate, and showcase a process of iterative model evaluation and improvement starting from a misspecified model and working toward a more appropriate representation of the data generating mechanism. For readers less familiar with dynamic systems modelling, skipping models 5 to 8 is suggested, while those more familiar may instead benefit from moving quickly over models 1 and 2, and noting that models 7 and 8 reflect relatively advanced topics that are not often used in psychological research at present.

Modelling Univariate Growth

Linear Change

Let us start by considering a hypothetical patient named Alex. Alex started going to couples therapy and has agreed to track their affect once every seven days (i.e. at equidistant time intervals), for 21 weeks, through a mobile app. What we want, is to conceptualize possible ways that Alex's affect develops during this time. A model in which nothing changes would be the simplest possible baseline. To begin with a minimally dynamic idea, however, we suppose that Alex's affect increases at a steady rate throughout the course of couples therapy.

Now we want to turn our conceptual idea of change into a statistical model so that we can fit it to data and obtain estimates of parameters. For a steady increase in affect, a linear regression model is an appropriate choice. With our linear regression model, we state that Alex's observed affect (y) changes at a steady rate (B) as a function of time (t) from an initial affect state (η_{t0}), along with deviations from this process. In this first example, we treat those deviations as random noise, unrelated to earlier or later time points – this could be considered as measurement error or some form of

variation that does not affect the overall trajectory of the process; later sections distinguish such observation-level error from system noise that changes the trajectory itself:

$$y(t) = \eta_{t0} + Bt + \epsilon(t)$$

At any time point t , Alex's observed affect equals their starting level (η_{t0}) plus a constant amount of growth per unit time (Bt), with some random measurement error ($\epsilon(t)$). The parameter B represents how much affect increases each week (with random deviations from this due to measurement error).

Next, we can obtain a better understanding of the growth process it implies (and whether it matches our conceptual idea) by generating data from the model (along with specific parameter values) and visualizing the resulting trajectory of affect.

Simulation

So let's write some code to generate data from a linear regression model representing a steady rate of change in affect and visualize its model-implied trajectory. The first code chunk below specifies a plotting function that we will re-use throughout the tutorial — details on this are beyond the scope of this tutorial, but essentially it contains a set of instructions for how to plot the data using the `ggplot2` package (Wickham, 2011), you should run this code before proceeding with further code chunks.

```
# Define reusable plotting functions
plot_trajectory <- function(data, x="Time", y="Affect",
  x_label = "Time (weeks)", y_label = "Affect",
  color_var = NULL, group_var = NULL, show_legend = FALSE, y2 = NULL){

  p <- ggplot(data, aes_string(x = x))
  if (!is.null(group_var)) {
    p <- p + aes_string(group = group_var)
  }
  # Handle single variable or couple plots
  if (is.null(y2)) {
    # Single variable plot
    p <- p + aes_string(y = y)
    p <- p + geom_point() + geom_line()
  } else {
    # Couple plot with two variables
    p <- p +
      geom_line(aes_string(y = y)) +
      geom_line(aes_string(y = y2), linetype = "dashed") +
      geom_point(aes_string(y = y))
  }
  if (!is.null(color_var)) {
    p <- p + aes_string(color = color_var)
    if (color_var %in% names(data) && is.numeric(data[[color_var]])) {
      p <- p + scale_color_gradient(low = "red", high = "blue")
    }
  }
  p <- p + theme_bw() + labs(x = x_label, y = y_label)
  if (!is.null(color_var) && is.null(y2)) {
    p <- p + labs(color = color_var)
  }
  if (!show_legend) {
    p <- p + theme(legend.position = "none")
  }
  return(p)
}
```

Moving to the actual simulation code, we first specify the number of time points we wish to generate data for (Time

<- 0:20). Second, we set the initial state of affect at the beginning of the growth process, i.e. the beginning of therapy (`t0Affect <- 5`). Third, we specify the linear regression model that will formalize our linear affect process for a single person, whose affect starts at an initial level (`t0Affect`) of 5 and increases steadily by 0.51 during every weekly observation (`Time*.51`). We use an arbitrary affect scale where larger values indicate more positive affect; the starting level, slope, and noise values are chosen to make the implied trajectory visible rather than to represent calibrated empirical values. Fourth, we add random noise to our model by sampling values from a normal distribution with a mean of 0 and a standard deviation of 1 (`rnorm(n=length(Time), mean = 0, sd = 1)`). These noise values are added to each affect value that is generated by our model. Fifth, we create a data frame to store the number of time points (i.e. weeks) and the affect values we generated for each time point. Finally, we use the created function `plot_trajectory` to plot the data using a lineplot to illustrate the model-implied affect process over the course of 21 weeks of therapy.

```
Time <- 0:20 # Generate time points for each measurement occasion
t0Affect <- 5 # Initial level of mood at week 0


# Specify linear model where affect increases by 0.51 each week
Affect <- t0Affect + Time*.51 +
  rnorm(n=length(Time), mean = 0, sd = 1)

# Create a data frame to store the generated data
data <- data.frame(Time = Time, Affect = Affect)

plot_trajectory(data)
```

Figure 1

Generated Linear Affect Trajectory for a Single Person.



Here we see our linear model provides a good representation of the idea that Alex's affect increases at a steady pace each time we observe them. Albeit with some random fluctuations, which in this simple model we attribute to unsystematic variation rather than to changes in the underlying affect process.

Dealing with Individual Differences in Affect Changes

By viewing Alex's model-implied growth process it strikes us that we have no idea whether their rate of change is slow or fast relative to other people who go to couples therapy. That is we may want to know if couples therapy leads to a different amount of change for different people, and if differences in people's rate of change are related to their initial level of affect. For instance, people who have a much lower initial level of affect than Alex may benefit more from couples therapy. To understand where Alex sits relative to other patients, we need to introduce individual differences into our model.

The equation for a linear model which allows people to vary in their initial affect level and rate of change, looks very similar to the linear regression model presented above but includes an i subscript to indicate that each person has their own initial level (η_{t0i}) and rate of change (B_{it}) of affect:

$$y_i(t) = \eta_{t0i} + B_{it} + \epsilon_i(t)$$

The subscript i indicates that each person i has their own parameters. Person i 's affect starts at their own initial level (η_{t0i}) and grows at their own rate (B_{it}). For example, Person 1 might start at affect level 3 and grow by 0.7 units per week, while Person 2 might start at affect level 6 and grow by 0.4 units per week. This allows us to model both within-person change and between-person differences.

This model is commonly termed a linear mixed-effects model (Bates et al., 2015), because it can include both fixed effects (parameters that are assumed to be the same across all subjects, often interpreted as population-level effects) and random effects (parameters that allow for individual variations, such as subject-specific intercepts η_{t0i} or slopes B_{it}).

Simulation

We now generate data from our linear mixed-effects model. The code below is used to generate data for multiple subjects (NSubjects) whose affect is measured 21 times, once per week for 21 weeks (times <- seq(from=0, to=20, by=1)). To generate a different initial state for each subject, we sample initial states from a normal distribution of initial affect levels with a mean of 5 and a standard deviation of 2 (t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)). The values here are arbitrary and can be tweaked to adjust the standard deviation of initial affect (sd), and the population average initial affect level (mean). To match our hypothetical scenario, where people's rate of growth in affect is related to their initial level, we construct each person's rate of change as a weighted combination of two standardized pieces: their initial affect level and an additional random slope component. The weight on initial affect is the desired intercept-slope correlation (initSlopeCor), and the weight on the random component keeps the total slope variance at 1 in expectation. We then shift and scale this standardized slope

using slopeMean and slopeSD, and check that the sampled correlation is close to the intended value.

```
# Generate data for multiple subjects with individual differences
NSubjects <- 20
times <- seq(from=0, to=20, by=1) #times when subjects are measured
Nobs <- length(times) #number of observations per subject
initSlopeCor <- -.8
#sample initial affect for each subject from normal distribution
t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)
#sample rate of change in affect for each subject
slopeMean <- .5
slopeSD <- .2
zInitialAffect <- as.numeric(scale(t0Affect)) #initial affect in SD units
randomSlopePart <- rnorm(n = NSubjects) #extra individual differences in slope
randomSlopePart <- as.numeric(scale(randomSlopePart)) #also in SD units
zTimeEffect <- #standardized slope with desired correlation to initial affect
  initSlopeCor * zInitialAffect + #correlated part
  sqrt(1 - initSlopeCor^2) * randomSlopePart #remaining random part
timeEffect <- slopeMean + slopeSD * zTimeEffect
cor(t0Affect, timeEffect) #print correlation
```

```
[1] -0.7660113
```

We can now create an empty data frame that can later hold the data we generate for each subject and each observation. To fill the data frame, we create two nested loops. The first loop (for(subi in 1:NSubjects)) allows us to go through each subject (subi) one by one. For each subject selected by the first loop, we initialize the second loop (for(obsi in 1:Nobs)) that allows us to go through each observation (obsi). Within this loop, we fill in each row in our data frame, which corresponds to an observation. For each observation, we generate information about a person's affect level, the current time point (week), and their subject identifier. To keep track of the current observation for each iteration of our loop, we initialize a row count at 0 (row <- 0) and then add 1 to this running row count (row <- row + 1). To reflect the imperfect measurement of affect, we add some random noise to the affect data that we generated, by sampling random values from a normal distribution with a mean of 0 and a standard deviation of 1 (data\$Affect <- data\$Affect + rnorm(n=nrow(data), mean = 0, sd = .1)). This measurement-error standard deviation is smaller than in the single-person plot above simply so that the between-person intercept-slope pattern is easily visible, not for substantive reasons.


```
#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA,NSubjects*Nobs),
  Time = rep(NA,NSubjects*Nobs),
  Affect = rep(NA,NSubjects*Nobs))

row <- 0 #initialize row counter to track current row
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1 #add 1 to row to reflect current row
    # current affect equals initial affect and effect of time:
    data$Affect[row] <- t0Affect[subi] +
      times[obsi] * timeEffect[subi]
    data$Time[row] <- times[obsi] #store time point for current row
    data$Subject[row] <- subi #store subject id for current row
  }
}
data$Affect <- data$Affect +
  rnorm(n=nrow(data), mean = 0, sd = .1) #measurement error
```

```
data$InitialAffect <- rep(t0Affect, each = Nobs)
plot_trajectory(data, color_var = "InitialAffect", group_var = "Subject", show_legend = TRUE)
```

Figure 2

Affect Trajectories for 20 Subjects with Individual Differences in Initial Level and Rate of Change.



From our plot, we can see that the correlated initial state and rate of change capture the idea that individuals who start with lower affect levels may improve faster over time.

Model Fitting

We can now fit our model to the data we generated and interpret how well it captures the patterns in our observations. To specify and fit our linear mixed-effects models, we will use the `lme4` package in R (explicitly including intercepts, although they would be included by default):

```
library(lme4)
lme_model <- lmer(Affect ~ # predict Affect
  1 + # a fixed effect for the intercept
  Time + # with a fixed effect of time
  # random intercept and random slope of time per subject:
  (1 + Time | Subject),
  data = data) # specify data to fit model to
```

After fitting the model to our data, we can ask for a summary of its parameter estimates.

```
# Summary of model output
summary(lme_model)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: Affect ~ 1 + Time + (1 + Time | Subject)
Data: data
```

```
REML criterion at convergence: -420.7
```

```
Scaled residuals:
    Min      1Q   Median      3Q      Max
-3.07762 -0.62579  0.05033  0.56896  2.79064
```

```
Random effects:
Groups   Name              Variance Std.Dev. Corr
Subject (Intercept)  4.310428  2.07616
          Time         0.033813  0.18388  -0.77
```

```
Residual          0.009913  0.09956
Number of obs: 420, groups: Subject, 20
Fixed effects:
              Estimate Std. Error t value
(Intercept)  5.22670    0.46434   11.26
Time         0.50128    0.04113   12.19

Correlation of Fixed Effects:
      (Intr)
Time -0.767
```

For this simple affect process, we can get a grasp on the model-implied trajectory through the numerical estimates offered by `lme4`'s output table. We can see that our linear mixed-effects model, recovers the parameter values from the data generating model well. We can find the population-level estimates for the intercept and slope values under "Fixed effects". The variance of the subject-specific deviations from the population average can be found under "Random effects". The correlation between the individual differences in the intercept and slope values is also under the "Random Effects" section (do not be misled by the "Correlation of Fixed Effects" heading, it does not reflect individual difference correlations). For more information about `lme4` and its functionalities we refer the reader to the package documentation (Bates et al., 2015) <https://cran.r-project.org/web/packages/lme4/lme4.pdf>.

Visualizing Predictions

Once our models become more complicated, a useful tool to understand their predictions is visualization. Visualization is a great way to build intuition about the process our models imply and allows us to detect sources of model misfit that may be difficult to identify if solely relying on numerical fit indices (see Gabry et al. (2019) for examples). For our `lme4` model, we can generate two plots showing: (1) the model's predictions (solid line) and its associated uncertainty (shaded ribbon) based only on the estimated model parameters for a single subject (subject 3); (2) predictions based on the model parameters and all data, (past, present, and future) of a single subject (subject 3). The code below uses the `bootMer` function to generate bootstrap (i.e., many) predictions from the model, and then uses the `ggplot2` package to visualize the predictions and uncertainty.

```
# Define a function to extract predictions
pred_fun <- function(model) {
  predict(model, data)
}

# Use bootMer to generate bootstrap (i.e., many) predictions
boot_preds <- bootMer(lme_model, FUN = pred_fun, nsim = 100)


# Extract predictions, calculate mean and confidence intervals:
# Create new data frame including original data:
newdata <- data.frame(data,
  pred = apply(boot_preds$t, 2, mean),
  lower = apply(boot_preds$t, 2, quantile, 0.025),
  upper = apply(boot_preds$t, 2, quantile, 0.975))

# Use ggplot2 to visualize the predictions and uncertainty
```

```
ggplot(newdata[newdata$Subject==3,], aes(x = Time)) +
  geom_line(aes(y = pred, color = "Predicted")) +
  geom_ribbon(aes(ymin = lower, ymax = upper,
    fill = "95% Confidence"), alpha = 0.2) +
  geom_point(aes(y = Affect, color = "Observed"), size = 2) +
  scale_color_manual(name = "",
    values = c("Predicted" = "blue", "Observed" = "red")) +
  scale_fill_manual(name = "",
    values = c("95% Confidence" = "blue")) +
  labs(y = "Affect", x = "Time") +
  theme_bw()+theme(legend.position = "bottom")
```

Figure 3

LME4 Model Predictions and Confidence Interval Based on Parameter Estimates Only, Subject 3.



The first prediction plot shows how well information solely from the estimated model parameters can reproduce our pattern of observations. The red dots show the observed affect values at each time point. The solid blue line shows the model implied expectation of affect, with the shaded blue area surrounding the line showing the variation / uncertainty about this trajectory. We can see that by solely relying on the model expectation, our model does a relatively poor job of predicting the actual data points. The solid line does not map well onto the observed trajectory of affect, and there is a large amount of uncertainty about the model-implied trajectory. This uncertainty is partly a consequence of partial pooling: before conditioning on the subject's own observations, the model relies mainly on the population distribution of trajectories rather than on the specific intercept and slope for this person.

```
# Define a function to extract predictions
preds <- predict(lme_model, se.fit = TRUE)


#create a new data frame including original data and predictions
newdata <- data.frame(data, preds,
  lower = preds$fit - 1.96 * preds$se.fit,
  upper = preds$fit + 1.96 * preds$se.fit)

# Use ggplot2 to visualize the predictions and uncertainty
ggplot(newdata[newdata$Subject==3,], #just visualise for subject 3
  aes(x = Time, y = fit)) +
  geom_line(aes(color = "Predicted")) +
  geom_ribbon(aes(ymin = lower, ymax = upper,
    fill = "95% Confidence"), alpha = 0.2) +
  geom_point(aes(y = Affect, color = "Observed"), size = 2) +
  scale_color_manual(name = "",
```

```
values = c("Predicted" = "blue", "Observed" = "red")) +
  scale_fill_manual(name = "",
    values = c("95% Confidence" = "blue")) +
  labs(y = "Affect", x = "Time") +
  theme_bw()+theme(legend.position = "bottom")
```

Figure 4

LME4 Model Predictions and 95% Confidence Interval for Subject 3 Conditional on All Observed Data.



In the second prediction plot, we visualize the model's predictions after it has learned from all the observations in our dataset. So, data from past, present, and future observations are used by the model to predict each observation. Here the model-implied trajectory does an excellent job at mapping onto our observations — having learnt the individual difference parameters needed for this subject specifically, predictions are far more accurate.

Nonlinear Change

By formalizing our linear growth process for Alex and other patients, we can inspect what our assumptions imply. In the simple simulated example above, the negative association between initial affect and rate of change implies that people with a lower initial level of affect can have a higher affect level at the end of the 21 weeks of therapy compared to people who started with higher affect. This means that if Alex entered therapy with more severe affect problems they would come out of therapy with a higher level of affect than if they entered with less severe affect problems. Moreover, if we take our model with a steady rate of growth seriously, then this would imply that Alex could endlessly increase their affect levels by continuing therapy. Both of these predictions seem unrealistic. So we have to rethink our model of affect (or interpret results from it **very cautiously**).

A more defensible prediction could be to expect that people's rate of growth slows down as their affect improves. This would mean that people with a low initial level of affect would increase faster at the start, relative to people with a higher initial level of affect, but their growth rate would slow down

as their affect improves. This is a conceptual shift from a deterministic trend indexed only by time to a state-dependent process: change depends on the current level of affect. It also has substantive implications. If a person's affect starts above the implied long-run level, the same model predicts movement downward rather than continued improvement (a more realistic scenario potentially relates individual differences in the start point to long-run level). In a statistical model, we can implement such a (nonlinear) growth process by allowing each person's slope to vary not just as a function of where they start, but also as a function of where they are *at every point in time*.

Our statistical model of this nonlinear affect process can still be represented using a regression model. We can achieve our nonlinear growth model by computing the observed level of affect at each time point ($\eta(t)$) as a function of affect at the previous time point ($A\eta_{t-1}$), plus a constant value (B) that is added to each observation:

$$\eta(t) = A\eta_{t-1} + B$$

At each time point t , affect depends on two components: (1) a fraction A of the previous affect level, and (2) a constant input B . The parameter A controls state dependence — if A is positive but less than 1 (e.g., $A = 0.8$), affect carries forward from the previous time point but at a reduced rate, creating a dampening effect — with only this state-dependency fraction, no matter the initial value of affect (whether positive or negative), the process would always converge towards zero. Including the constant input B allows the process to converge towards a different, non-zero value (which is not exactly B , but a function of both A and B). In case of an A between 0 and 1, and a B value greater than zero, the process will converge towards a positive value — whether the process starts above this value or below. Essentially, this creates nonlinear growth: when affect is low, the total of the state-dependency and constant input is larger (faster growth). As affect increases, growth slows down, as the loss from the state-dependency fraction becomes similar to the gain from the constant input B .

With this equation we can emulate a process where the rate of growth in affect dampens throughout the course of a therapy session. Our code will do the heavy lifting of computation and allow us to visualize the model-implied trajectory that our equation creates. But we can also build some intuition by solving this equation for a small number of time points.

So, let's consider 3 time points T0, T1, and T2. Then let's set the autoregressive coefficient (A) at 0.8, which means that 0.8 of the previous observation will be added to the current observation. Let's also add a constant of 2, as our continuous intercept (B) which is added to every observation.

Then, we can start up our process at T0 by setting an initial level of affect to 5 ($\eta(0) = 5$). At T1, we can see that $0.8 * 5 + 2$

leads to a current level of affect $\eta(1) = 6$. By applying our equation to the next time point, T2, we can see that $0.8 * 6 + 2$ leads to current level of affect equal to 6.8. So from T0 to T1 our affect level grew by 1 unit, but from T1 to T2 it only grew by 0.8. Thus, the rate of growth is dampening as consequence of a change in the prior level of affect.


Simulation

Below we provide code that does this computation for all of our 21 weeks, using the equation we just discussed. First we need to set the degree to which the current level of affect is dependent on the prior level of affect (i.e. state dependence, $A <- .8$). Then we add a constant value which is added to each observation generating a consistent rate of growth ($B <- 2$). By adding the state dependence term with the continuous intercept, we generate a nonlinear growth trajectory. We then set an initial affect level which is used to initialize our process ($t0Affect <- 5$). An if-else statement is used so that the first observation sets the initial level of affect ($\text{if}(i==1) \text{Affect}[i] <- t0Affect$) and each subsequent observation has an affect value that is generated by our model equation ($\text{else } \text{Affect}[i] <- A * \text{Affect}[i-1] + B$).

```
times <- seq(from=0, to=20, by=1) #time points of measurement
A <- .8 #autoregressive coefficient / state dependence
t0Affect <- 5
B <- 2 #continuous intercept
Affect <- rep(NA, length(times)) #create empty affect vector
for(i in 1:length(times)){ #for each measurement occasion
  #if first time point, set to initial affect:
  if(i==1) Affect[i] <- t0Affect
  #otherwise, use autoregression and continuous intercept:
  else Affect[i] <- A*Affect[i-1] + B
}
plot_trajectory(data.frame(Affect=Affect, Time=times))
```

Figure 5

Nonlinear Trajectory from Discrete-Time Autoregressive Model.



Our figure shows a nonlinear increase in affect over time. This is a better representation of a reality where therapy can-

not make you infinitely happy, but it may help you recover from a low point.

Modelling Univariate Dynamic Systems

A Conceptual Shift to Continuously Evolving Processes

Our model now provides a more realistic depiction of Alex's growth in affect from week to week. But what happens in between each of our weekly observations? One possibility is that affect stays flat and suddenly jumps during each of our weekly observations. That is, change only happens when we are looking. If we try to imagine such a discretely changing affect process, it would look like a staircase of affect. The alternative possibility is that people's affect changes continuously across time, forming a coherent trajectory. The rate of change that we observe from week to week, is then an aggregate of all the moment-to-moment changes that happen between our weekly observations. If the second option is more akin to our expectation of reality, then a continuous-time framework can better represent the process. This becomes particularly important when modelling multiple interacting processes or when dealing with data collected at varying time intervals (For details on such issues we refer the reader to [Driver, 2025](#); [Kuiper & Ryan, 2018](#); [Voelkle et al., 2012](#)).

Differential Equations — The Mathematics of Continuous Time

Differential equations are a natural way to represent processes that are assumed to evolve continuously. A researcher does not need a deep technical understanding of differential equations to build useful intuition: the key shift is from modelling the value of affect at the next observation based on the current observation, to modelling how fast affect is changing at the present moment based on its current value. One way to soften this shift is to start with average change over a finite interval. As the interval Δt becomes smaller, this average rate approaches the instantaneous rate of change, written as the derivative $\frac{d\eta}{dt}$. The differential equation then describes what determines that momentary rate of change.

Modelling a growth process in continuous time will be no different than in discrete time, in that we will use symbols to represent different factors and their relationships which combine to form our growth process. To understand the equation that generates our hypothesized affect process we can again visualize the model-implied process using simulations.

We can model an affect process that changes at a steady rate in continuous time, akin to our first example where we used a linear regression model — but this time using a differential equation. The major difference here is that our outcome variable (left-hand side of equation) is the *rate of change*, rather than the *level*, of affect — level was the outcome in our discrete time (regression) representation. This rate of change

in affect (η) at a given moment in time t is denoted by the derivative $\frac{d\eta}{dt}$. We can think of the rate of change at a given moment in time as what we would get if we were to glance at a speedometer in our car. The velocity that we would see, would tell us how fast our current position is changing at a particular moment in time. Since our model is linear, the rate of change at a given moment is given by a constant value B — there is no variation in the rate of change (or slope), it remains constant over time. Thus, the differential equation for a linear model is:

$$\frac{d\eta}{dt} = B$$

The notation $\frac{d\eta}{dt}$ reads as “the rate of change of η with respect to time t ” — think of it as affect’s “velocity” at any instant. If $B = 0.51$, then at every moment, affect is increasing at a constant rate of 0.51 units per week. This is like a car traveling at a constant speed, always travelling 100 km/h. Similarly, no matter Alex’s current affect level, it’s always increasing by 0.51 units per week.

For readers who remember basic calculus, this is the derivative of a linear function: if $\eta(t) = \eta(0) + Bt$, then the derivative with respect to time is B . The continuous-time equation therefore describes the same linear idea, but in terms of the instantaneous slope rather than the level at a particular observation.

If we want to pick up where we left off and represent our nonlinear growth model in continuous time, we can add a term that allows the rate of growth at a given time point to depend on the current state of affect (i.e. a state dependency):

$$\frac{d\eta}{dt} = A\eta(t) + B$$

Now the rate of change (velocity) depends on where affect currently is, as well as on the constant input B . The term $A\eta(t)$ means the rate of change is proportional to current affect level. If A is negative (e.g., $A = -0.2$), higher affect leads to slower growth (or even decline), creating a self-regulating system — just as we saw in our discrete time model when the autoregression coefficient A was positive but less than 1. Here in the continuous-time case A is still a regression coefficient, but the outcome variable is now the **rate of change in affect**, rather than the **level of affect at a new time point**. The term B is still a constant upward push.

Simulation

To generate data for this continuous growth process in R, we can approximate the solution to the differential equation using a ‘Euler-Maruyama’ method ([Higham, 2001](#)) (or technically in this case without system noise, we would call this just the ‘Euler method’). This method works by breaking continuous time into small time slices (steps) and updating the

state of our growth process at each step based on the model's dynamics.

The differential equation defines an instantaneous rate of change, but simulation requires us to approximate the trajectory over finite steps. To generate data, we first define values for our model parameters. This includes: (1) the degree to which the current level of affect influences its own rate of change (A, continuous time state-dependence). (2) A constant input to the growth process (B, continuous intercept), and an initial level of affect (t0Affect). Then we create an empty vector to store the affect values given by our for-loop which implements the 'Euler-Maruyama' method (`Affect <- rep(NA, Nobs)`). The for-loop sets the initial level of affect for the first observation (`i=1`) to initialize the process (`if(i==1) Affect[i] <- t0Affect`). For each subsequent observation (`i > 1`), we compute the current level of affect by summing the *rate of change* at the prior time point (`dAffect = A*Affect[i-1] + B`) with level of affect at the prior time point (`Affect[i-1]`). That is, `Affect[i] <- Affect[i-1] + dAffect * 1`. To add the size of the desired time step into the equation, we multiply the rate of change by the size of the time step, in this case `dAffect` is multiplied by 1, meaning that the rate of change is equivalent to the rate of change between each observation.


```
times <- seq(from=0, to=20, by=1) #measurement times
Nobs <- length(times) # number of observations per subject
A <- -.2 # continuous time state dependence
t0Affect <- 3 # set initial level of affect
B <- 2 # continuous intercept
Affect <- rep(NA, Nobs) # create empty affect vector

for(i in 1:Nobs){ # for each time point
  # if first time point, set to initial affect
  if(i==1) Affect[i] <- t0Affect
  else{ # otherwise compute slope of affect at earlier time point
    dAffect <- A*Affect[i-1] + B
    # then update affect using slope and time step
    Affect[i] <- Affect[i-1] + dAffect * 1
  }
}

plot_trajectory(data.frame(Affect=Affect, Time=times))
```

Figure 6

Generated Affect Trajectory from Continuous-Time Model.



This is a somewhat crude approximation however — it assumes the rate of change is constant over each time interval of 1, when in fact our model specifies that the rate of change varies whenever affect changes, which is continuously. We use this rough approximation first for intuition, then refine it. To better approximate a continuous time solution, we can take smaller steps in time. But more steps means that the necessary computation will be more intensive. So let's start with a small number of intermediary steps, say 10 steps in the time between our observations, to better approximate the continuous time process `Nsteps <- 10`. Now, we introduce another loop, within our previous for-loop, which splits the interval between each observation (e.g. week) into 10 smaller steps (`istep`). The first line within this loop represents the differential equation which gives us the rate of change at the prior time point (`dAffect = A*Affect + B`). The line below it, updates the current level of affect based on the estimated rate of change at the prior time point multiplied by 1/10 (i.e. 1/Nsteps). The effect of this multiplication is to give us the rate of change over the smaller step of time we chose. In essence, this scales the rate of change from the full unit of time (e.g., week), given by the preceding equation, to one-tenth, allowing us to better approximate the underlying continuous growth process.

```
times <- seq(from=0, to=20, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
A <- -.2 #continuous time state dependence
t0Affect <- 3
B <- 2 #continuous intercept
Affect <- rep(NA, Nobs) #create empty affect vector
Nsteps <- 10 #number of steps between each observation

for(i in 1:Nobs){ #for each time point
  #if first time point, set to initial affect
  if(i==1) Affect[i] <- t0Affect
  else{ #compute new affect state by sequence of small steps
    #initialise with state at previous time point
    currentAffect <- Affect[i-1]
    for(stepi in 1:Nsteps){ #take Nsteps for each observation
      #compute slope of affect at earlier time point
      dAffect <- A*currentAffect + B
      #update state using slope and time step
      currentAffect <- currentAffect + dAffect * 1/Nsteps
    }
    Affect[i] <- currentAffect
  }
}
```

Nonlinear Change with Random Fluctuations (Continuous Time)

By using our previous model to predict Alex's affect trajectory, we would assume that their affect would change in a smooth deterministic manner. However, real-world affect processes are often subject to unmodelled influences that lead a person's affect to fluctuate in meaningful and persistent ways — perhaps Alex wakes up with a headache one day, or has a particularly nice conversation with a friend. These may not alter their long term trajectory, but can induce an effect both immediately and over some length of time. Such fluctuations are distinct from variation we might attribute to

measurement error where we typically assume that imperfections in our measurements are random and independent at each occasion of measurement.

We can model random fluctuations in the affect process by adding a system noise term (sometimes called process noise or latent variance) to our differential equation. In our model so far, affect depends only on its current level and a fixed intercept, but in reality many other variables also matter: sleep, minor hassles, chance encounters, and countless other influences we do not measure or include explicitly. System noise stands in for the combined effect of these unmodeled variables: it allows the true affect state to be pushed up or down by influences outside the equation. Unlike measurement error, which affects only the observation, system noise changes the underlying affect level itself, and those shifts persist because the dynamics carry them forward over time. With our current equation, the rate of change in affect $\frac{d\eta}{dt}$ is purely deterministic, meaning that how fast affect changes at a given moment is entirely determined by the terms in the right-hand side of the differential equation ($A\eta(t) + B$), with no allowance for such unmodeled influences.

To model the impact of random influences on Alex's affect trajectory, we want to add a term that introduces a random “kick” to the rate of change at each moment. Meaning that aside from the deterministic effects there will be some random ups and downs in Alex's growth rate. To add this “kick” to Alex's rate of change we can introduce a noise term that is sampled from a normal distribution. This turns our (ordinary) differential equation into a stochastic differential equation (Kloeden & Platen, 1992), with stochastic denoting the random noise component of our continuous process.

That is, we model our continuous time process as evolving over arbitrarily small steps in time. At each very small step in time, our noise term adds a portion of randomness to the rate of change. The amount that it adds is scaled by the system noise term G , defining the volatility of the process. Thus, our rate of change in affect $\frac{d\eta(t)}{dt}$ is a combination of some deterministic effect ($A\eta(t) + B$) and some random effect ($GdW(t)$), over each interval dt . Because of the stochastic nature of the random noise term, it is not possible to analytically define its rate of change at any point in time. Thus, stochastic differential equations are often written a bit differently, with the dt term moved to the right-hand side of the equation, only applied to the deterministic portion:

$$d\eta = (A\eta(t) + B)dt + GdW(t)$$

This equation now has two components to change: (1) *Deterministic change* ($A\eta(t) + B$) dt : the predictable trend based on current affect level, and (2) *Random fluctuations* $GdW(t)$: unpredictable “shocks” that can push affect up or down randomly at any moment. This reflects the idea that while there is some predictability to change in affect over time, there are also unpredictable changes that can occur at any moment,

which can persist over time

Simulation

Adding system noise. To add system noise to our continuous time model, we will update the data generating code using the Euler method discussed in the previous section. The system noise term will be added to the affect level that is computed for each of our chosen time steps ($istep$). Thus, when we update a person's current affect state at each small time step we add some noise. This noise is sampled from a normal distribution with a mean of 0 and a standard deviation of 1 divided by $Nsteps$ (this time 100 steps). Before being added to the function that generates the current affect state, each draw of random noise is multiplied by G (the system noise coefficient) to scale the amount of system noise. Thus, the higher the G value the more noisy the affect process becomes.

Individual differences. To generate data for multiple individuals we have to make some further tweaks to our code. First, for each participant we sample their initial affect level from a normal distribution with mean 5 and standard deviation of 2 ($t0Affect \leftarrow rnorm(n = NSubjects, mean = 5, sd = 2)$). Second, we add an empty data frame which includes a greater number of rows to fill. Third, we add a loop to repeat the computation of the affect process for each subject ($for(subi \text{ in } 1:NSubjects)$).

```
NSubjects <- 20
times <- seq(from=0, to=20, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.1 #continuous time state dependence
B <- 1 #continuous intercept
G <- 0.2 #system noise coefficient


#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA, NSubjects*Nobs),
  Time = rep(NA, NSubjects*Nobs),
  Affect = rep(NA, NSubjects*Nobs))

Nsteps <- 100 #steps between each observation
row <- 0 #init current row tracker
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1
    #if first time point, set to initial affect
    if(obsi==1) Affect <- t0Affect[subi]
    if(obsi>1){ #else compute new state via small steps in time
      for(stepi in 1:Nsteps){ #every step, do...
        dAffect <- A*Affect + B #compute affect slope
        Affect <- Affect + #new state = old state +
          dAffect * 1/Nsteps + #deterministic change +
          G * rnorm(n=1, mean=0, sd=sqrt(1/Nsteps)) #random change
      }
    }
    data$Affect[row] <- Affect #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Subject[row] <- subi #input subject data
  }
}

#add measurement error
data$Affect <- data$Affect+rnorm(n=nrow(data), mean = 0, sd = .05)
#plot trajectories
plot_trajectory(data, color_var = "as.factor(Subject)")
```

Figure 7

Affect Trajectories with System Noise and Individual Differences.



Now we have a somewhat more realistic depiction of affect dynamics, with fluctuations around the smooth, state-dependent increase in affect over time. This model captures the idea that affect processes are subject to both predictable trends and random fluctuations, reflecting some of the complexity of real-world affect dynamics. Note also, that here individuals vary in their initial levels of affect, and the rate of change varies accordingly, such that therapy does not just help everyone at the same rate, and nor does it lead to unrealistic long term projections as we saw in the linear case — there is a limit.

ctsem: A Flexible Tool for Dynamic Systems Modelling

To fit a continuous-time model with state-dependence and system noise to data we now shift from hand-coded generative simulation to formal estimation in a ‘state-space’ framework (Durbin & Koopman, 2012; Harvey, 1989; Kalman, 1960). This shift adds a layer of abstraction. Instead of treating the observed affect score as the process itself, we distinguish a latent dynamic process from the observations that measure it. The latent process evolves continuously according to the dynamic model, while the measurement model links that latent process to the observed affect scores at the times they were measured.

The examples in this tutorial use simple single-indicator measurement models for pedagogical clarity. Thus, the tutorial emphasizes the dynamic state-space part of continuous-time structural equation modelling more than full multiple-indicator latent-variable measurement modelling. Multiple indicators, factor loadings, and more elaborate measurement structures can be specified in this framework, but they are not the focus of this tutorial – for a tutorial focused on R and SEM multiple-indicator measurement modelling see Rosseel and Vidal (2025).

The R package *ctsem* (Driver et al., 2017; Driver & Voelkle, 2018a) provides a software framework for specifying

and fitting continuous-time (CT) and discrete-time (DT) dynamic models. It enables researchers to model the behavior and evolution of processes over time, and its combination of features can offer some unique advantages for analyzing longitudinal data.

Software Installation

Before starting, ensure your system is ready to use *ctsem*. Follow these steps:

1. **Install R and RStudio:** Download and install R (<https://cran.r-project.org/>) and RStudio (<https://posit.co/products/open-source/rstudio/>) or alternative interfaces.
2. **Configure c++ for model compilation with Stan:** *ctsem* relies on Stan (Carpenter et al., 2017) for log likelihood calculations. Most of the models we will fit do not require compilation, but some of the more complex models we will fit do. Ensure you have a compatible C++ compiler installed:
 - **Windows:** Install RTools (<https://cran.r-project.org/bin/windows/Rtools/>).
 - **Mac:** Install Xcode command line tools by running `xcode-select --install` in the terminal.
 - **Linux:** Install g++ and other build tools using your package manager (e.g., `sudo apt install build-essential`).
3. **Install Required Packages:**

```
install.packages("ctsem", dependencies=TRUE)
```

Model Fitting

To fit the data we just generated with *ctsem*, we first need to specify our model using the `ctModel` function. The code below is annotated for convenience, for more information about the specification of a continuous-time model using the `ctModel` function please refer to the *ctsem* manual (<https://cran.r-project.org/web/packages/ctsem/ctsem.pdf>) (Driver et al., 2025). Here we will provide a quick description of some of the links between the model equation and the software specification.

Our *ctsem* model can be thought of as linking a latent (unobserved) ‘true’ affect process to observed (or ‘manifest’) variables, which are part of our data. Observed variables are treated as indicators of the underlying latent process, and are linked to the latent states via a measurement model. For the running example, this means that Alex’s affect is assumed to evolve continuously between observations, while the observed ratings are (potentially) noisy snapshots of that underlying process.

In this case our latent dynamic model is given by the stochastic differential equation:

$$d\eta = (A\eta(t) + B)dt + GdW(t)$$

This model is quite general, and can be formulated as very large matrices (when e.g., there are many latent processes interacting with each other and/or many observed variables). As such `ctsem` expects matrices as inputs for each of the different model components, but also contains some shortcuts to make it easier to specify the model. Model components can be either fixed to specific values (input a numeric value) or estimated from the data (input a string with the name of the parameter to estimate). When we move on to matrices later, each matrix can be a mixture of fixed and estimated values. Here is how we specify our current dynamic model using the `ctsem` syntax:

1. The **DRIFT** argument specifies the state dependence A , and contains a freely estimated parameter we call `stateDependence`, making the rate of growth depend on the current level of affect.
2. The **CINT** argument specifies the continuous intercept (B). We call the continuous intercept parameter B and only estimate a fixed effect (one value for all subjects). To switch off random effects (on by default for initial values and intercept parameters) we need to add `FALSE` into the 3rd ‘slot’ of the parameter, done by adding vertical bars after the parameter name, i.e. `B || FALSE`. We do not need to modify the 2nd slot, so it is left empty.
3. The **DIFFUSION** argument specifies the system noise term (G), and contains a freely estimated parameter `systemNoise`, allowing the model to capture random fluctuations around the model-implied growth trajectory.
4. The initial value of our latent process is treated as a random variable with its own mean and variance parameters: $\eta_{i0} = \mu_{0i} + \epsilon_{0i}$, where $\epsilon_{0i} \sim \mathcal{N}(0, \sigma^2)$. The **TOMEANS** argument, is used to specify the mean parameter for the initial level of affect, which we call `t0Affect`. By adding `TRUE` behind `t0Affect ||` we are allowing for random effects (between-subject variance) around the population estimate of the initial level of affect in our sample — this would have been enabled by default in any case.

The latent process model is linked to our observations using a (linear) measurement model:

$$\text{Affect}(t) = \Lambda\eta(t) + \tau + \epsilon(t)$$

Where $\text{Affect}(t)$ is the observed level of affect. Λ is a matrix of regression weights used to link the latent process to our observations, η is the latent level of affect, τ reflects a constant that can be used to shift the relationship between observations and a latent variable, and ϵ represents measurement error, and is distributed as a multivariate normal distribution with mean 0 and (usually diagonal) covariance matrix Θ .

1. The **LAMBDA** argument specifies the Λ matrix. In our case it is a 1×1 matrix with the value 1, meaning that the latent Affect (η) is directly and equally reflected in the observed Affect without any scaling.

2. The **MANIFESTMEANS** argument specifies the measurement intercept (τ), this is a constant that is added to the measurement model to shift the relationship between observations and a latent variable. In this case it is set to zero, implying that there is no systematic offset in the measurement process).
3. The **MANIFESTVAR** argument specifies the measurement error variance in terms of standard deviations (ϵ), which we call `residualSD`.

```
m_basic <- ctModel( #define the ctsem model
# Specify features of the data / model type
manifestNames = "Affect", #names of observed variables in data
latentNames = "Affect", #names for latent processes
time = 'Time', #name of time column in dataset
id = 'Subject', #name of subject column in dataset
type='ct', #use continuous-time (dt for discrete-time)
# Specify model components
MANIFESTVAR = 'residualSD', #SD of measurement error
LAMBDA = matrix(1,nrow=1,ncol=1), #factor loading matrix
MANIFESTMEANS=0, #set measurement intercept / offset to 0
CINT='B||FALSE', #continuous intercept with *no* random effects
TOMEANS='t0Affect||TRUE', #initial affect + random effects
DRIFT = 'stateDependence', #state dependence
DIFFUSION = 'systemNoise') #system noise
```

Checking and understanding model specification can be much easier when looking at the full, expanded model equations in matrix form — we can generate a LaTeX formatted representation of the equations we just specified by running the `ctModelLatex()` function. To do this we need to install the `tinytex` package, if we have not already done so. Sometimes there are difficulties getting LaTeX compilation to work on various systems. For those who have difficulties, generally installing the R package `tinytex` via `install.packages('tinytex')`, then running `tinytex::install_tinytex()`. In case of errors along the way, restarting R / Rstudio and clearing the workspace can help. The model output is shown in Figure 8, which brings together the model elements introduced above. The subject-level distribution describes individual differences in initial affect. The dynamic model describes how the latent affect state changes through drift, continuous intercept, and diffusion elements. The measurement model describes how the latent state becomes the observed affect score plus measurement error.

We can now fit the model and ask for a summary of the parameter estimates.

```
f_basic <- ctStanFit(dataalong = data, ctstanmodel = m_basic)
```

Figure 8

Matrix Representation of the Basic Univariate ctsem Model.

Subject parameter distribution: $\underbrace{[t0Affect_i]}_{\phi(i)} \sim \text{tform} \{N([raw_t0Affect], [rawPCov_1.1])\}$

Initial latent state: $\underbrace{[Affect] (t_0)}_{\eta(t_0)} \sim N \left(\underbrace{[t0Affect]}_{TOMEANS}, \underbrace{UcorSDtoCov \{ [1e-06] \}}_{Q^{*}_{t0} \text{ TOVAR}} \right)$

Deterministic change: $\underbrace{d[Affect] (t)}_{d\eta(t)} = \left(\underbrace{[stateDependence]}_{\mathbf{A} \text{ DRIFT}} \underbrace{[Affect] (t)}_{\eta(t)} + \underbrace{[B]}_{\mathbf{b} \text{ CINT}} \right) dt +$

Random change: $\underbrace{UcorSDtoChol \{ [systemNoise] \}}_{\mathbf{G} \text{ DIFFUSION}} d \underbrace{[W_1] (t)}_{d\mathbf{W}(t)}$

Observations: $\underbrace{[Affect] (t)}_{\mathbf{Y}(t)} = \underbrace{[1]}_{\mathbf{A} \text{ LAMBDA}} \underbrace{[Affect] (t)}_{\eta(t)} + \underbrace{[0]}_{\mathbf{\tau} \text{ MANIFESTMEANS}} +$

Observation noise: $\underbrace{UcorSDtoChol \{ [residualSD] \}}_{\mathbf{\Theta} \text{ MANIFESTVAR}} \underbrace{[e_1] (t)}_{\epsilon(t)}$

System noise distribution per time step: $\Delta[W_{j \in [1,1]}](t-u) \sim N(0, t-u)$

Observation noise distribution: $[\epsilon_{j \in [1,1]}](t) \sim N(0, 1)$

```
summary(f_basic, parmatrices=FALSE) #some output disabled
```

```
$residCovStd
  Affect
Affect 0.082

$resiCovStdNote
[1] "Standardised covariance of residuals"

$popstd
      mean    sd  2.5%   50%  97.5%
t0Affect 1.9233 0.311 1.381 1.9028 2.5837

$popmeans
      mean    sd  2.5%   50%  97.5%
t0Affect 5.3411 0.4204 4.5229 5.3462 6.1526
stateDependence -0.0933 0.0071 -0.1074 -0.0931 -0.0797
systemNoise 0.2075 0.0134 0.1821 0.2073 0.2325
residualSD 0.0649 0.0222 0.0319 0.0620 0.1176
B 0.9460 0.0573 0.8302 0.9469 1.0524

$popNote
[1] "covariance pars in sd / unconstrained cor form, see $parmatrices for cor/cov"

$loglik
[1] 8.340722

$npars
```

```
[1] 6
```

```
$aic
[1] -4.681443
```

```
$logposterior
[1] 8.340722
```

```
$parmatNote
[1] "For additional summary matrices, use argument: parmatrices = TRUE"
```

The summary output provides several key sections for interpreting model fit and parameters:

- **\$residCovStd**: Standardized residual covariance, characterising the amount of unexplained variation in the data after the models' predictions have been accounted for.
- **-PARMATRICES**: Parameter matrices for the model, suppressed in shown output.
- **\$rawpopcorr**: Correlations between random effects in the model – only a single random effect in this case, so not shown.

- `$popstd`: Population standard deviations for random effects in the model
- `$popmeans`: Population-level parameter estimates for fixed effects model parameters
- `$loglik`: Log-likelihood value for model comparison
- `$aic`: Akaike Information Criterion for model selection

From our summary table, we can see that the estimated coefficients are not exactly the same as the true values, this is due to sample variation in the data. But the model captures the general trend of affect dynamics over time. The estimated state dependence coefficient is around -0.1, the continuous intercept (B) is approximately 1, and the system noise standard deviation is around 0.2. The random effects standard deviation for initial affect is approximately 2, and the residual standard deviation is around 0.05.


Visualizing Predictions

Using `ctsem` we can easily plot the model predictions and associated uncertainty. Our predictions can be conditional on all, none, or some of the individual subjects' observed data. First let's look at the predictions based solely on the parameter estimates, for an arbitrary subject, number 6 (multiple subjects can be specified in the `subjects` argument).

```
ctKalman(fit= f_basic, plot = TRUE, subjects = 6, removeObs = TRUE)
```

Figure 9

ctsem Predictions for Subject 6 Based on Parameter Estimates Only.




Just like in the case of our simple linear model, we see that the model does a poor job at predicting the observations (red dots) and there is a large degree of uncertainty (shaded red ribbon) around the model-implied trajectory (solid red line). Now let's allow the model to learn from every 5th observation to supplement the predictions it makes using the model estimates alone – of course, when the model is being estimated, it uses every data point for prediction, but visuals such as this

can help provide the intuition for how the model works (note the updates to the prediction trajectory when new data is observed) and also provide useful checks on the model, in case predictions diverge substantially.

```
ctKalman(fit= f_basic, plot = TRUE, subjects = 6, removeObs = 5)
```

Figure 10

ctsem Predictions for Subject 6 Conditional on Parameter Estimates and Every 5th Observation.




We can see that the model already does a better job at predicting the observations, and the uncertainty in its predicted trajectory is reduced (though growing larger again towards the end of the time period before a new observation is fed into the model). Lastly, let us include all past and present observations (the default) and see how our model performs.

```
ctKalman(fit= f_basic, plot = TRUE, subjects = 6, removeObs = FALSE)
```

Figure 11

ctsem Predictions for Subject 6 Conditional on All Data.



Now the model does a fairly good job at predicting the observations, but there is still some uncertainty in its predictions.

Time-Dependent Predictors: Therapy Sessions as Impulse Effects (Continuous Time)

So far, we have modeled individual differences using time-independent predictors — variables that remain constant across time for each person (like personality traits or relationship duration). However, many influences on psychological processes vary within persons over time. For example, therapy sessions, stressful life events, medication doses, or social interactions can all have time-varying effects on affect. These are captured using *time-dependent predictors* (TDpreds).

In *ctsem*, time-dependent predictors operate (by default) as *impulse effects*: they provide an instantaneous “shock” to the latent process at the moment they occur. Critically, the effect of this impulse then feeds through the system over time, just as with any other inputs (such as the random fluctuations we’ve already introduced). In self-regulating systems, where inputs tend to dissipate back towards equilibrium, such an impulse effect will also gradually decay back toward the equilibrium. This provides a natural model for transient events — such as a prompt, stressor, medication dose, social interaction, or experimental manipulation — whose effect is expected to be strongest when it occurs and then dissipate through the system dynamics.

Consider Alex, who attends weekly therapy sessions during their 20-week treatment period. Each session provides a temporary boost to their affect. The longer term effect of the therapy is already accounted for by the growth trajectory we have already specified in earlier examples. Between sessions, the short term effect of therapy fades away and their affect gradually returns toward their general growth trajectory. The key insight is that we do not need to separately model the “decay” of the therapy effect — it emerges naturally from the state dependence dynamics we have already specified.

This is a stylized input example rather than a substantive theory of psychotherapy. Real therapy effects may be cumulative, delayed, heterogeneous, mediated by unobserved mechanisms, or confounded with events that determine when sessions occur. The present example is used to show how an event at a known time can be represented as an input to a dynamic process.

Mathematically, if a therapy session occurs at time t_k with effect size M , the latent state receives an instantaneous impulse:

$$\eta(t_k^+) = \eta(t_k^-) + M$$

where $\eta(t_k^-)$ is the state just before the session and $\eta(t_k^+)$ is the state just after.

Simulation

To simulate data with therapy sessions as impulse effects, we create a time-dependent predictor *Therapy* that equals 1 at observations when a therapy session occurs and 0 other-

wise. For simplicity, we schedule therapy sessions at regular intervals (every 3 time units).

```
# Generate data for multiple subjects with therapy sessions
NSubjects <- 20
times <- seq(from=0, to=20, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.2 #continuous time state dependence
B <- rnorm(n = NSubjects, mean = 1, sd = .3) #CINT
G <- .3 #system noise coefficient
M <- 1.5 #therapy impulse effect (positive = boost to affect)

# Define therapy sessions (every 3 time units, starting at time 3)
therapyTimes <- seq(from=3, to=max(times), by=3)

#create empty data.frame to fill step by step
data <- data.frame(Subject = rep(NA, NSubjects*Nobs),
  Time = rep(NA, NSubjects*Nobs),
  Affect = rep(NA, NSubjects*Nobs),
  Therapy = rep(NA, NSubjects*Nobs))

Nsteps <- 10 # steps between observations

row <- 0 #initialize row counter
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){
    row <- row + 1
    currentTime <- times[obsi]
    therapyToday <- ifelse(currentTime %in% therapyTimes, 1, 0)


    if(obsi == 1) Affect <- t0Affect[subi]
    if(obsi > 1){ #update state via small steps in time
      for(stepi in 1:Nsteps){
        dAffect <- A * Affect + B[subi] #deterministic slope
        #update state using slope and time step
        Affect <- Affect + dAffect * 1/Nsteps +
          G * rnorm(n=1, mean=0, sd=sqrt(1/Nsteps)) #system noise
      }
    }
    #add impulse if therapy session (dummy variable)
    Affect <- Affect + M * therapyToday
    data$Affect[row] <- Affect #add affect to data
    data$Time[row] <- currentTime #add time to data
    data$Subject[row] <- subi #add subject to data
    data$Therapy[row] <- therapyToday #add therapy to data
  }
}

#add measurement error (after creating process!)
data$Affect <- data$Affect+rnorm(n=nrow(data), mean = 0, sd = .1)

plot_trajectory(data[data$Subject %in% 1:4, ], #plot
  color_var = "as.factor(Subject)" + #then overlay therapy
  geom_vline(xintercept=therapyTimes, linetype="dashed", alpha=.5))
```

Figure 12

Affect Trajectories for Four Subjects with Therapy Sessions as Impulse Effects (Vertical Lines).



The dashed vertical lines indicate therapy sessions. Notice how affect tends to spike upward at these times and then follows the general growth trajectory we specified in earlier examples.

Model Fitting

To fit this model in `ctsem`, we add the Therapy variable as a time-dependent predictor using the `TDpredNames` argument. By default, `ctsem` estimates the effect of time-dependent predictors on each latent process, here we explicitly name the parameter `td_Affect_Therapy`. The estimated effect represents the instantaneous impulse — the immediate impulse to the latent state when the predictor equals 1. We will not estimate individual differences in this effect, but one could imagine such differences being of interest if going further with the model.

```
m_therapy <- ctModel(
  manifestNames = "Affect",
  latentNames = "Affect",
  TDpredNames = "Therapy", #time-dependent predictor
  TDPREDEFFECT = c('td_Affect_Therapy'),
  time = 'Time',
  id = 'Subject',
  type = 'ct',
  MANIFESTVAR = 'residualSD',
  LAMBDA = matrix(1, nrow=1, ncol=1),
  MANIFESTMEANS = 0,
  CINT = 'B||TRUE', #continuous intercept with random effects
  TOMEANS = 't0Affect||TRUE',
  DRIFT = 'stateDependence',
  DIFFUSION = 'systemNoise')
```

```
f_therapy <- ctStanFit(dataalong = data, ctstanmodel = m_therapy)
```

```
s <- summary(f_therapy)
print(s$popmeans)
```

	mean	sd	2.5%	50%	97.5%
t0Affect	4.5534	0.4003	3.8385	4.5469	5.3225
stateDependence	-0.1849	0.0131	-0.2111	-0.1848	-0.1590
systemNoise	0.2843	0.0221	0.2427	0.2840	0.3289
residualSD	0.1127	0.0276	0.0704	0.1103	0.1780
td_Affect_Therapy	1.4742	0.0353	1.4073	1.4726	1.5446
B	0.9320	0.1129	0.7036	0.9320	1.1584

The `td_Affect_Therapy` parameter represents the estimated impulse effect of therapy sessions. A positive value indicates that therapy boosts affect. Our simulated value was $M = 1.5$, so we expect an estimate close to this.

Modelling Multivariate Systems

Couple Dynamics with Self and Partner Effects (Continuous Time)

Until now, we have looked at Alex's affect in isolation. But Alex is going to couples therapy, which makes it very likely that how their partner, Robin, feels has a pronounced impact on how Alex feels. In the same way, Alex's affect is also likely to have a reciprocal impact on Robin's affect. This creates an interdependence between Alex and Robin's affect dynamics — a fundamental concept in dyadic data analysis where the behavior or state of one partner influences and is influenced by the other partner, creating a complex system of mutual dependencies. The same modelling structure is useful in a broad array of contexts: for example, symptoms may influence one another, physiological and behavioral processes may be coupled, and parent-child or therapist-client systems may involve reciprocal dynamics. The statistical challenges of dyadic interdependence (Kenny et al., 2006) are particularly relevant here as we transition from modelling individual dynamics to modelling coupled systems where cross-partner effects and shared environmental influences must be explicitly accounted for. We can model this interdependence in two ways:

1. We can introduce a *cross-effect state dependence term*, which reflects how Alex's affect at a given moment can have a direct effect on how Robin's affect changes at that moment. This is essentially the state dependence we modeled within a person, but now the state (affect level) that drives the rate of change is also coming from another person. We can also introduce another cross-effect state-dependence term, capturing the influence of Robin's current state of affect on Alex's rate of change at that moment.
2. *Correlation in system noise*, random life events that cause fluctuations in Alex's affect trajectory are also likely to contemporaneously impact Robin's affect (and vice versa). We can model factors that simultaneously shift our couple's rate of change, by allowing a random noise term for each partner, but allowing some correlation in the noise — capturing the idea that Alex and Robin each experience unique influences, but also some shared influences. Such shared influences could include a household stressor, a conflict, a shared positive event, the same measurement context, or another unobserved environmental input that affects both partners.

The revised model which includes Alex and Robin's affect dynamics along with their interdependence can be written just as earlier (the dimensions of the model components will differ however, and for simplicity we ignore the time dependent predictor for now), or as two equations, one for each partner:

$$d\eta_1 = (A_{11}\eta_1(t) + A_{12}\eta_2(t) + B_1)dt + G_{11}dW_1(t) + G_{12}dW_2(t)$$

$$d\eta_2 = (A_{21}\eta_1(t) + A_{22}\eta_2(t) + B_2)dt + G_{21}dW_1(t) + G_{22}dW_2(t)$$

We have two equations, one for each partner's affect. The equation for $d\eta_1$ (rate of change of Alex's affect), implies that Alex's rate of change depends on (1) their own current affect level ($A_{11}\eta_1(t)$ — the auto-effect), (2) their partner Robin's current affect level ($A_{12}\eta_2(t)$ — the cross-effect from partner), (3) a constant input (B_1), and (4) random noise that may affect both partners ($G_{11}dW_1(t) + G_{12}dW_2(t)$). In general, we would not be able to estimate G_{12} distinctly from G_{21} and vice versa, as they are correlated, so when it comes to estimation we will estimate a single parameter to represent the correlation between them, as well as the standard deviations of the noise for each partner. The equation for $d\eta_2$, rate of change in Robin's affect, implies that similarly, Robin's rate of change depends on their own affect ($A_{22}\eta_2(t)$), Alex's affect ($A_{21}\eta_1(t)$), their constant (B_2), and random noise that may affect both partners ($G_{21}dW_1(t) + G_{22}dW_2(t)$). The cross-effects (A_{12} and A_{21}) capture how partners influence each other's rate of change. If A_{12} is positive, when Robin's affect is high, Alex's affect is positively influenced — it either grows faster or declines slower.

Simulation

We will generate data for 20 couples. Each of the 40 people will have their own initial level of affect, but all dynamics and growth terms will be fixed across individuals for now.

Couple dynamics. To generate data for each couple, we can build on the previous model of affect that we used to represent the affect of one person. The first thing we need to add, is an equation to represent the affect process of the additional partner. Then we need to add our cross-effect and correlated-noise terms to the (differential) equations. In the code chunk below, we can see that we now have two equations to approximate two continuous affect processes, which are differentiated using `Affect1` (for partner 1) and `Affect2` (for partner 2). Apart from duplicating the equation for one person, each equation has *two* new elements to capture interdependency in affect dynamics. (1) The change in affect for each person depends on their partner's level of affect at the same time point. The coefficient of this cross-effect is denoted by `Across` (e.g. `dAffect1 <- A*Affect1 + Across * Affect2 + B`). (2) We allow the random fluctuations in the couple's affect process to covary. We do this by generating correlated noise according to a specified `DIFFUSIONcov` covariance matrix,

where the diagonal elements specify the variance for each process and the off-diagonal elements specify the covariance between processes.

```
NSubjects <- 20
times <- seq(from=0, to=40, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect1 <- rnorm(n = NSubjects, mean = 5, sd = 2)
t0Affect2 <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.2 #continuous time state dependence
Across <- .1 #cross-effect state dependence
B <- 1 #continuous intercept
#covariance matrix for system noise
DIFFUSIONcov <- matrix(c(1, .05, .05, .1), nrow=2, ncol=2)

#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA, NSubjects*Nobs),
  Time = rep(NA, NSubjects*Nobs),
  Affect1 = rep(NA, NSubjects*Nobs),
  Affect2 = rep(NA, NSubjects*Nobs)) #now with two individuals

Nsteps <- 100 #steps between observations

row <- 0 # initialize row counter
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1 #increment row counter
    if(obsi==1){ #set initial affect
      Affect1 <- t0Affect1[subi]
      Affect2 <- t0Affect2[subi]
    }
    if(obsi>1){ #update state via small steps in time
      for(stepi in 1:Nsteps){ #for each step
        # compute deterministic slopes of affect
        dAffect1 <- A*Affect1 + Across * Affect2 + B
        dAffect2 <- A*Affect2 + Across * Affect1 + B

        # generate correlated noise
        systemNoise <- MASS::mvrnorm(n=1, mu=c(0,0),
          Sigma=DIFFUSIONcov/Nsteps) #scaled by stepsize

        #update state using slopes and time step + noise
        Affect1 <- Affect1 + dAffect1/Nsteps + systemNoise[1]
        Affect2 <- Affect2 + dAffect2/Nsteps + systemNoise[2]
      }
    }
    data$Affect1[row] <- Affect1 #input affect data
    data$Affect2[row] <- Affect2 #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Subject[row] <- subi #input subject data
  }
}

#add measurement error
data$Affect1 <- data$Affect1+rnorm(n=nrow(data), mean = 0, sd=.05)
data$Affect2 <- data$Affect2+rnorm(n=nrow(data), mean = 0, sd=.05)

data$Couple <- factor(data$Subject) #factor subject for plotting
plot_trajectory(data[data$Subject %in% c(1,2),], y="Affect1",
  show_legend=TRUE, y2="Affect2", color_var="Couple")
```

Figure 13

Affect Trajectories for Both Partners in One Couple.




Figure 13 shows the affect trajectories for two couples. Each partner in the couple is clearly distinct, but couples show similarities in comparison to other couples, due to both shared influences and effects from one partner on the other.

Model Fitting

First, the model syntax for `ctsem` for the couple model:

```
m_couple <- ctModel( #define the ctsem model
  manifestNames = c("Affect1", "Affect2"), #observed variables
  latentNames = c("Affect1", "Affect2"), #names of latent processes
  time = 'Time', #name of time column in dataset
  id = 'Subject', #name of subject column in dataset
  type='ct', #use continuous time (dt for discrete-time)
  MANIFESTVAR = c( #covariance matrix of measurement error
    'residualSD1', 0, #sd on diagonal, 0 correlation assumed.
    0, 'residualSD2'),
  LAMBDA = diag(1,2), #diagonal factor loading matrix (identity)
  MANIFESTMEANS=c(0,0), #zero measurement intercept / offset
  CINT=c('B1|F', 'B2|F'), #continuous intercept no random effects
  TOMEANS=c('t0Affect1||TRUE',
    't0Affect2||TRUE'), #initial affect with random effects
  DRIFT = c(
    'auto1', 'cross12', #auto effect partner 1, cross-effect 2 to 1
    'cross21', 'auto2'), #cross-effect 1 to 2, auto effect for 2
  DIFFUSION = c(
    'systemNoise1', 0, #system noise for 1, 0 in upper triangle
    'systemNoiseCor', 'systemNoise2')) #corr, and sd for 2
```

To understand the `ctsem` syntax for our multivariate model (one latent process and one observed variable for each partner) it helps to see the model equation in expanded matrix form to more easily link it with the more compact syntax of `ctsem`. We use `ctModelLatex(m_couple)` for this purpose, shown in Figure 14. Here, both partners' equations are combined into matrix equations.

The first matrix equations in Figure 14 show the subject parameter distribution, confirming that we have two parameters in our model that vary across individuals, `t0Affect1` and `t0Affect2`, and these are assumed normally distributed according to some estimated mean and covariance, *possibly* with some transformation applied after this distribution to ensure parameters remain within necessary bounds (primarily

relevant when e.g., correlation or standard deviation parameters are estimated, see the `ctsem` manual for more details on this).

The second matrix equation in Figure 14 refers to the distribution of the initial latent state **within-subject** – because we have individual differences in the initial states, these are not necessary (not uniquely estimable), so fixed to a small variance and determined based on the subject specific parameters `t0Affect1` and `t0Affect2`. The `TOMEANS` argument is used to specify these. By adding `TRUE` behind `t0Affect1||` and `t0Affect2||` we are allowing for random effects (between-subject variance) around the population estimate of the initial level of affect in our sample.

Next we look at the dynamic model component of Figure 14, made up of the deterministic change and random change lines. On the left, the vector $(dAffect_1, dAffect_2)$ represents the rates of change for both partners. On the right, the drift matrix shows how each partner's current affect influences each partner's rate of change. The diagonal elements (A_1, A_2) are auto-effects (how each person's own affect influences their own change), while off-diagonal elements (A_{12}, A_{21}) are cross-effects (how one partner affects the other's change). The `DRIFT` argument to `ctModel` specifies this matrix. The `CINT` (continuous intercept) vector contains the constants (B_1, B_2) that are steady inputs for each partner, and is specified by the `CINT` argument to `ctModel`. Because in this case we only estimate a fixed effect (one value for all subjects), we switch off random effects (on by default for initial values and intercept parameters, i.e., `TOMEANS`, `CINT`, `MANIFESTMEANS`) by adding `FALSE` behind `B||`, i.e. `B||FALSE`.

Turning to the random change in Figure 14, it is created by a complex function (`UcorSDtoChol`) that takes the *DIFFUSION* matrix as input, where the diagonal parameters are standard deviations and the lower-triangle off-diagonal parameters represent 'unconstrained correlations' between the noise terms. The `ctModel` function takes *DIFFUSION* as an argument to specify this matrix of standard deviations and unconstrained correlations. Our standard deviation parameters are `systemNoise1` and `systemNoise2` and the correlation parameter is called `systemNoiseCor`. These unconstrained correlation parameters are unfortunately somewhat complex to interpret directly, but are necessary to allow for continuous-moderation and random-effects. For interpretation, obtaining the *DIFFUSION* covariance from the `summary()` or `ctStanContinuousPars()` output is recommended, possibly with a conversion to a correlation matrix via the `cov2cor()` function.

The latent process model is linked to our observations using a (linear) measurement model, given by the observation equation in Figure 14. This equation also comprises a deterministic component and a random component, and links what we actually observe ($Affect_1(t), Affect_2(t)$) to the underlying

Figure 14

Matrix Equations for the Bivariate Couple ctsem Model.

$$\begin{aligned}
&\text{Subject parameter distribution: } \underbrace{\begin{bmatrix} t0Affect1_i \\ t0Affect2_i \end{bmatrix}}_{\phi(i)} \sim \text{tform} \left\{ N \left(\begin{bmatrix} \text{raw.t0Affect1} \\ \text{raw.t0Affect2} \end{bmatrix}, \begin{bmatrix} \text{rawPCov.1.1} & \text{rawPCov.2.1} \\ \text{rawPCov.2.1} & \text{rawPCov.2.2} \end{bmatrix} \right) \right\} \\
\\
&\text{Initial latent state: } \underbrace{\begin{bmatrix} Affect1 \\ Affect2 \end{bmatrix}}_{\eta(t_0)}(t_0) \sim N \left(\underbrace{\begin{bmatrix} t0Affect1 \\ t0Affect2 \end{bmatrix}}_{\text{TOMEANS}}, \underbrace{UcorSDtoCov \left\{ \begin{bmatrix} 1e-06 & 0 \\ 0 & 1e-06 \end{bmatrix} \right\}}_{\mathbf{Q}_{t_0}^* \text{ TOVAR}} \right) \\
\\
&\text{Deterministic change: } \underbrace{d \begin{bmatrix} Affect1 \\ Affect2 \end{bmatrix}}_{d\eta(t)}(t) = \underbrace{\begin{bmatrix} \text{auto1} & \text{cross12} \\ \text{cross21} & \text{auto2} \end{bmatrix}}_{\mathbf{A} \text{ DRIFT}} \underbrace{\begin{bmatrix} Affect1 \\ Affect2 \end{bmatrix}}_{\eta(t)}(t) + \underbrace{\begin{bmatrix} B1 \\ B2 \end{bmatrix}}_{\mathbf{b} \text{ CINT}} dt + \\
\\
&\text{Random change: } \underbrace{UcorSDtoChol \left\{ \begin{bmatrix} \text{systemNoise1} & 0 \\ \text{systemNoiseCor} & \text{systemNoise2} \end{bmatrix} \right\}}_{\mathbf{G} \text{ DIFFUSION}} d \underbrace{\begin{bmatrix} W_1 \\ W_2 \end{bmatrix}}_{\mathbf{W}(t)}(t) \\
\\
&\text{Observations: } \underbrace{\begin{bmatrix} Affect1 \\ Affect2 \end{bmatrix}}_{\mathbf{Y}(t)}(t) = \underbrace{\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}}_{\mathbf{A} \text{ LAMBDA}} \underbrace{\begin{bmatrix} Affect1 \\ Affect2 \end{bmatrix}}_{\eta(t)}(t) + \underbrace{\begin{bmatrix} 0 \\ 0 \end{bmatrix}}_{\mathbf{\tau} \text{ MANIFESTMEANS}} + \\
\\
&\text{Observation noise: } \underbrace{UcorSDtoChol \left\{ \begin{bmatrix} \text{residualSD1} & 0 \\ 0 & \text{residualSD2} \end{bmatrix} \right\}}_{\mathbf{\Theta} \text{ MANIFESTVAR}} \underbrace{\begin{bmatrix} \epsilon_1 \\ \epsilon_2 \end{bmatrix}}_{\epsilon(t)}(t) \\
\\
&\text{System noise distribution per time step: } \Delta[W_{j \in [1,2]}](t-u) \sim N(0, t-u) \quad \text{Observation noise distribution: } [\epsilon_{j \in [1,2]}](t) \sim N(0, 1)
\end{aligned}$$

Note: *UcorSDtoChol* converts lower tri matrix of standard deviations and unconstrained correlations to Cholesky factor, *UcorSDtoCov* = transposed cross product of *UcorSDtoChol*, to give covariance, See Driver & Voelkle (2018) p11.
Individual specific notation (subscript i) only shown for subject parameter distribution – pop. means shown elsewhere.

latent processes ($Affect_1(t)$, $Affect_2(t)$) that represent ‘true’ affect.

The deterministic component of the measurement model is given by $\Lambda\eta(t) + \tau$, where Λ (LAMBDA argument to *ctModel*) is a matrix of regression weights used to link the latent process to our observations, and is here set to the identity matrix (1s on diagonal, 0s off-diagonal), meaning each observed affect measure directly reflects its corresponding latent affect with a 1-to-1 relationship (no scaling or transformation). τ (MANIFESTMEANS argument to *ctModel*) reflects a constant that can be used to shift the relationship between observations and a latent variable, but here is a zero vector, meaning there’s no systematic bias or offset in measurements (This becomes particularly useful when estimating factor models, but can also be estimated instead of CINT, such that latent processes converge towards zero and MANIFESTMEANS offsets this an appropriate amount for the data).

The random component of the measurement model (ob-

servation noise) is given by $\epsilon(t)$, where ϵ is a multivariate normal distribution with mean 0 and (usually diagonal) covariance matrix Θ . The Θ matrix (MANIFESTVAR argument to *ctModel*) contains measurement error standard deviations (residualSD1, residualSD2) — these scale the disturbance terms ϵ to capture the idea that our observations are imperfect indicators of the true latent affect state, but that there is no correlation between the imperfections of each — the off-diagonal elements are here set to 0 (but could be estimated if desired, as with the DIFFUSION correlations).

An important note when specifying the system matrices in *ctsem*, is that they can be explicitly created using the *matrix()* function in R, but can also be passed in as a vector (except for the LAMBDA, factor-loading matrix) using the *c()* function, and this vector will be read in **row-wise** in contrast to typical R convention (this makes model specification more intuitive to read as it aligns with the matrix representation).

Fit and Summarize ctsem Model

Then we fit our model and could, as earlier, use `summary` to view estimated parameters. However, since it is difficult to get a feeling for the nonlinear model-implied trajectory of a system composed of multiple interdependent process through numerical values, we will focus on a visual representation of our system instead for the time being, with the exception of code showing the correlation in system noise.

```
f_couple <- ctStanFit(datalong = data, ctstanmodel = m_couple)
parmatrices <- ctStanContinuousPars(f_couple) #estimated matrices
print(cov2cor(parmatrices$DIFFUSIONcov)) #system noise correlation
```

```
          Affect1  Affect2
Affect1  1.0000000  0.1110163
Affect2  0.1110163  1.0000000
```


Visualize Predictions

Just like before, we can look at how well the model predicts our observations based on the estimated model parameters plus every 5th data point for our 3rd dyad (subject 3). When we visualize predictions from our model, now including multivariate dynamics, state dependence, and random fluctuations, we see a very different picture to the earlier linear models. Our model is able to capture a reality where there is a complex interplay between the nonlinear affect dynamics of two individuals. The uncertainty in the predictions reflects two sources: The first, is that there are inherently unpredictable fluctuations in affect due to factors we have not observed or modeled. So our model is inevitably imperfect. The second, is that each measurement of affect is likely an imperfect indicator of the underlying affect state, so we can't be sure of the true state of the system even at the moment where we measure it (via e.g., a survey on a smartphone). We could visualise the latent process predictions and uncertainty by changing the `kalmanvec` argument to `'etaprior'` instead of `'yprior'` (which shows predictions for observations at each moment).

```
ctKalman(fit = f_couple, plot=TRUE, subjects = 3,
         kalmanvec=c('y','yprior'), removeObs = 5)
```

Figure 15

ctsem Predictions for Both Partners in Dyad 3 Conditional on Every 5th Observation.




Visualize Dynamics — Independent Shocks

We can also use visualizations to build intuition about the model-implied dynamics (i.e. auto-effect and cross-effect state dependencies). To do this, we can create a plot that tracks the results of an imaginary experiment. In this experiment, we take both Alex and Robin, and place them in a completely controlled setting, where no external factors can impact their affect. That is, their affect remains at baseline unless we do something about it. Second, we intervene on Alex's affect to increase it by 1 at time 0. Third, we plot how Alex's and Robin's affect would be expected to change over 10 weeks following the shock to Alex, as a function of their individual auto-effects and the cross-effects they have on each other. For example, if we increase Alex's affect by 1 at time 0, Alex's affect is expected to still be approximately 0.40 above baseline after 5 weeks (red solid line, `Affect1.Affect1` in the plot at time interval = 5) because of self-dependency and the bidirectional cross-effects between Alex and Robin. Instead, Robin's affect doesn't jump initially, but rises more gradually in response to the higher affect of Alex — it is expected to be 0.20 by week 5 (green solid line, `Affect2.Affect1` in the plot at time interval = 5). We can also visualize what happens if we increase Robin's affect by 1 *instead of Alex's* by viewing the blue and purple solid lines. This plot is known as an impulse response function plot (Lütkepohl, 2005).

```
ctStanDiscretePars(f_couple, plot=T)
```

Figure 16


Impulse Response Functions Assuming Independent Shocks*Visualize Dynamics — Correlated Shocks*

Our hypothetical experiment with independent shocks, may have low generalizability to real world dynamic systems. This is because often a shock that happens to one partner also affects the other partner, and even though we may model this correlation, we have only been able to estimate model parameters in this correlated context — the impulse response plot (and direct interpretation of state dependence coefficients) may not generalise well to the hypothetical intervention scenario depicted above (see Driver, 2023, for more on the combined interpretation of system noise and dynamics).

One way to consider this issue is to create a plot where we consider the estimated correlation in the system noise, which tells us to what extent Alex and Robin's fluctuations are related, and incorporate this into the initial impulse — depicting essentially a more 'typical' impulse for the estimated system. This essentially means that when Alex experiences a random fluctuation in affect, Robin is also likely to experience a fluctuation, depending on the estimated correlation. Instead of plotting an impulse response function based solely on an independent 1-unit shock, we modify the impulse response function to reflect that a shock to one partner comes with a correlated fluctuation in the other partner's affect. For example, if the system noise correlation is estimated at 0.25, then when Alex's affect increases by 1 unit, Robin's affect is expected to change by approximately 0.25 units as a result of that shared noise component. With `ctsem` we can simulate and plot the impulse response function, such that when one partner is shocked by 1 unit the response in the other partner is scaled by the system noise correlation. This gives a more realistic picture of the dynamic interplay between the partners.

```
ctStanDiscretePars(f_couple, plot=T, observational=TRUE)
```

Figure 17

Impulse Response Function Assuming Correlated Shocks

From this plot we see that when Alex's affect increases by 1 at T_0 , Robin's affect is expected to contemporaneously fluctuate in the same direction as well (positive system noise correlation). The positive cross-effect coefficients also show that there is a positive bidirectional relationship between the two partners. Such that when one partner's affect increases we can expect that to also cause a subsequent increase in the other partner's affect. Here, we can be sure that this reflects a causal effect because we generated the data with such an effect. But in a real-world scenario, we would need to consider the possibility of confounding variables that could explain this estimated relationship. The system noise correlation can account for confounders that change very quickly (compared to the speed of change of affect), but is unlikely to sufficiently account for confounders that change at a similar speed to affect. For instance, a sudden loud noise that startles both partners will lead to a very brief fluctuation in affect for both partners, which could be reasonably reflected in the system noise. However, a prolonged period of gloomy weather might cause a gradual systematic shift in the affect of both partners that can be misinterpreted as evidence that one partner's affect directly impacts the other's affect. Such slowly changing confounders are not accounted for in this model at present, but allowing for stable differences (in e.g., the continuous intercept) is one way that very slow / stable confounders can be accounted for — more on this later.

Individual Differences in System Dynamics

In the above models, we have generally assumed that all individuals have the same system dynamics. In reality, people may have highly heterogeneous system dynamics. This is because there are likely to be an immense number of factors that contribute to what we might consider as the individual's dynamic system. Factors that change slowly or not at all, with respect to our observed time window, can reasonably be treated as stable individual differences. Such individual differences may exist in the magnitude of random fluctuations,

with some people having more stable affect systems while others have a more volatile affect. Alternatively, in the case of couples, some dyads may have more interdependent affect dynamics with stronger coupling, while other dyads are more independent.

Couple Dynamics with Individual Differences and Time-Independent Moderation (Continuous Time)

Stable moderators are useful whenever people or systems may differ in their dynamic parameters because of relatively enduring characteristics. In other psychological applications, these moderators might be diagnostic group, age, trait affect, baseline stress, treatment condition, relationship length, socioeconomic context, or a stable environmental exposure. A key question is whether the moderator plausibly changes the dynamic process or measurement model itself, rather than merely predicting average levels.

For our case, we are going to consider two sources of individual differences. First, we are going to assume that individuals have differences in their long-term, post therapy baseline affect. This will be done by allowing individual differences in the continuous intercept. Allowing for this makes our model a continuous-time analogue to the popular random intercept cross-lagged-panel model (Hamaker et al., 2015). Second, we will assume that couples who have been together longer have a stronger influence on each others' affect. This will be incorporated by allowing for individual differences in the strength of our cross-effect state-dependence terms, and by allowing the size of each couple's cross-effect to depend on the average time they spend together.

Simulation

The code chunk to run this simulation, builds on the code we used earlier. To generate a continuous intercept for each person we sample from a normal distribution, $B1 \leftarrow \text{rnorm}(n = \text{NSubjects}, \text{mean} = 2, \text{sd} = .3)$. We assume that people with a higher affect are more likely to get in a relationship, thus we add a common value sampled from a normal distribution to the value for each person in a dyad. That is, $B_{\text{common}} \leftarrow \text{rnorm}(n = \text{NSubjects}, \text{mean} = 1, \text{sd} = .2)$ is added to the sampled values of $B1$ and $B2$ ($B1 \leftarrow B_{\text{common}} + \text{rnorm}(n = \text{NSubjects}, \text{mean} = 2, \text{sd} = .3)$).

To add individual differences in the continuous intercept parameter, we index the parameter to denote that it is specific to the denoted subject. So $B1$ becomes $B1[\text{subi}]$ where subi refers to the subject identifier.

A similar process is used to generate data for our heterogeneous cross effects. We sample a cross-effect for each couple from a normal distribution and then add the time the couple has spent together, scaled by a coefficient that determines the degree of influence time spent together has on each couple's cross effects $\text{Across} \leftarrow \text{rnorm}(n = \text{NSubjects}, \text{mean} =$

$.1, \text{sd} = .05) + .05 * \text{TogetherZ}$. The cross-effect parameter is also indexed in our data generating loop to generate a subject-specific parameter, e.g. $\text{Across}[\text{subi}] * \text{Affect2}$.

```
NSubjects <- 20
times <- seq(from=0, to=40, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect1 <- rnorm(n = NSubjects, mean = 5, sd = 2)
t0Affect2 <- rnorm(n = NSubjects, mean = 5, sd = 2)
Together <- runif(n=NSubjects, min=0, max=20) #years together
TogetherZ <- scale(Together) #standardise time together
A <- -.4 #continuous time state dependence
#cross-effect state dependence -- affected by time together.
Across <- rnorm(n=NSubjects, mean = .1, sd=.05) + .05*TogetherZ

# generate continuous intercepts for each individual,
# assuming high affect individuals may partner with high affect
# common continuous intercept variance:
Bcommon <- rnorm(n = NSubjects, mean = 1, sd = .2)
B1 <- Bcommon + rnorm(n = NSubjects, mean = 2, sd = .3) #partner 1
B2 <- Bcommon + rnorm(n = NSubjects, mean = 2, sd = .3) #partner 2

# covariance matrix for system noise
DIFFUSIONcov <- matrix(c(.1, .05, .05, .1), nrow=2, ncol=2)

#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA, NSubjects*Nobs),
  Time = rep(NA, NSubjects*Nobs),
  Affect1 = rep(NA, NSubjects*Nobs),
  Affect2 = rep(NA, NSubjects*Nobs)) #now affect for two individuals

Nsteps <- 100 # steps between each observation
row <- 0 #init current row tracker
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1
    if(obsi==1){ # if first time point, set to initial affect
      Affect1 <- t0Affect1[subi]
      Affect2 <- t0Affect2[subi]
    }
    if(obsi>1){ # else update state via small steps in time
      for(stepi in 1:Nsteps){ # for each step

        #compute deterministic slopes
        dAffect1 <- A*Affect1 + Across[subi] * Affect2 + B1[subi]
        dAffect2 <- A*Affect2 + Across[subi] * Affect1 + B2[subi]

        # generate correlated noise scaled by stepsize
        systemNoise <- MASS::mvrnorm(n=1, mu=c(0,0),
          Sigma=DIFFUSIONcov/Nsteps)


        #update states using slope and time step, add noise
        Affect1<- Affect1 + dAffect1 * 1/Nsteps+systemNoise[1]
        Affect2<- Affect2 + dAffect2 * 1/Nsteps+systemNoise[2]
      }
    }
    data$Affect1[row] <- Affect1 #input affect data
    data$Affect2[row] <- Affect2 #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Subject[row] <- subi #input subject data
    data$TogetherZ[row] <- TogetherZ[subi] #input time together
  }
}

# add measurement error
data$Affect1 <- data$Affect1+rnorm(n=nrow(data),mean=0,sd=.05)
data$Affect2 <- data$Affect2+rnorm(n=nrow(data),mean=0,sd=.05)

plot_trajectory(data[data$Subject==1,],y="Affect1",y2="Affect2")
```

Figure 18

Affect Trajectories for Both Partners in One Couple with Time-Together Moderation.



Model Fitting

These individual differences can be accommodated in *ctsem* using both ‘fixed’ and ‘random’ effect approaches. In the fixed effect approach, we rely on observed covariates (time-independent predictors in *ctsem*) to moderate the system parameters and explain the individual differences. In *ctsem* these are restricted to linear effects at present (a complication here is that parameters such as correlations or standard deviations have non-linear transforms applied afterwards, to ensure they cannot go out of bounds — see the *ctsem* manual for more detail). For alternative forms, the usual approach of including higher order terms (e.g., age and age squared) is possible. Creative use of *ctsem*’s modelling flexibility could also be used to estimate other nonlinear effects.

In the random effect approach, we assume that the individual differences are at least partly due to unobserved factors, and estimate the variance and correlations in these factors, relying purely on the observed data (without covariates) to estimate where each individual’s parameter(s) sit with respect to their distribution of subject-specific deviations from the average population estimate.

By combining fixed and random effects, we try to improve the estimated model for each individual by allowing individuals to have unique system dynamics, but not ‘completely unique’ — we still rely to some extent on how other individuals behave (for more on continuous-time hierarchical modelling see Driver & Voelkle, 2018a). In *ctsem* we can either decide to automatically specify random effects (preferable in most cases), or manually add random effects via the specification of additional latent processes (for details on the manual method see the Driver & Tomasik, 2023). For our example, we will leave the specification of the random effects to the *ctsem* back-end, which estimates all random-effects as correlated.

So, let’s go over the new syntax in our *ctsem*. In terms

of data that we fit our model on, we add the time couples spent together (TogetherZ) as a stable source of individual differences (time-independent predictor). This is done using the argument `TIpredNames = c('TogetherZ')`. By default when covariates are included in *ctsem* they moderate all parameters of the system. Thus, we turn the default moderation off using `tipredDefault = FALSE`. Now if we want our time-independent predictor to moderate a specific parameter, we need to explicitly specify the moderation. The moderation of a cross-effect (e.g. `cross21`) by TogetherZ can be specified within the DRIFT matrix by inserting the name of the predictor to the fifth slot (slots are denoted by the pipe `|` operator, where each `|` defines 1 slot). Thus, we write `cross21||TRUE||TogetherZ`. This would estimate the `cross21` parameter, allow it to vary as a linear function of TogetherZ, and allow for random effects (done by adding `TRUE` to the third slot, `cross21||TRUE`) to account for any individual differences that could not be explained by TogetherZ. It is important to note that moderators in *ctsem* can dramatically influence the interpretation of parameters and lead to confusion if care is not taken — usually centering, and possibly scaling, the moderator(s) is a good idea. Since we also want random effects for the continuous intercept of each partner we specify these using `CINT=c('B1||TRUE','B2||TRUE')`. All the other arguments are explained in prior sections of this tutorial, or can be found in the *ctsem* manual (<https://cran.r-project.org/package=ctsem/vignettes/hierarchicalmanual.pdf>).

```
# Fit continuous time structural equation model
m_indiff <- ctModel( #define the ctsem model
  tipredDefault = FALSE, #moderation disabled unless specified
  TIpredNames = c('TogetherZ'), #time independent predictors
  manifestNames = c("Affect1","Affect2"), #observed variables
  latentNames = c("Affect1","Affect2"), #names of latent processes
  time = 'Time', #name of time column in dataset
  id = 'Subject', #name of subject column in dataset
  type='ct', #use continuous time
  MANIFESTVAR = c('residualSD1',0,
    0, 'residualSD2'), #sd's of the residual / measurement error
  LAMBDA = diag(1,2), #factor loading matrix
  MANIFESTMEANS=c(0,0), #no measurement intercept / offset
  CINT=c('B1||TRUE','B2||TRUE'), #cint with random effects
  TOMEANS=c('t0Affect1||TRUE','t0Affect2||TRUE'), #initial affect
  DRIFT = c( #auto and cross effects matrix, cross effects with...
    'auto1', 'cross21||TRUE||TogetherZ', #random effects and...
    'cross21||TRUE||TogetherZ','auto2' ), #moderation
  DIFFUSION = c( #system noise matrix
    'systemNoise1', 0, #sd for subj 1, 0 in upper triangle
    'systemNoiseCor', 'systemNoise2')) #correlation, sd for subj 2

f_indiff <- ctStanFit(datalong = data, ctstanmodel = m_indiff)

s = summary(f_indiff) #store summary of the fit
```

Now if we look at the `$tipreds` (time-independent predictors) section of the summary, we have an estimated value for the effect of time together on the cross-effect state dependence. We can see that the cross-effect state dependence increases by approximately .05 for each increase of 1 in standardized time together. Thus, partner’s who spent more time together have more tightly coupled affect.

```
print(s$tipreds)
```

	mean	sd	2.5%	50%	97.5%	z
tip_TogetherZ_cross21	0.0427	0.0115	0.0193	0.0429	0.0661	3.7052

The `$rawpopcorr` section of the summary shows the estimated correlations between the random effects of the model — the initial states and the continuous intercepts. We can see that the correlation between the random effects for the continuous intercepts (the `B2__B1` parameter) is approximately the true correlation we used when generating the data.

```
print(round(s$rawpopcorr,2))
```

	mean	sd	2.5%	50%	97.5%	z
t0Affect2__t0Affect1	0.11	0.20	-0.32	0.12	0.50	0.56
cross21__t0Affect1	-0.23	0.29	-0.70	-0.26	0.38	-0.79
B1__t0Affect1	0.56	0.26	-0.13	0.62	0.87	2.12
B2__t0Affect1	0.21	0.27	-0.37	0.23	0.67	0.76
cross21__t0Affect2	0.10	0.31	-0.52	0.12	0.64	0.33
B1__t0Affect2	-0.02	0.29	-0.55	-0.02	0.52	-0.06
B2__t0Affect2	-0.23	0.26	-0.68	-0.24	0.33	-0.86
B1__cross21	-0.20	0.33	-0.74	-0.24	0.45	-0.60
B2__cross21	-0.29	0.31	-0.79	-0.31	0.36	-0.94
B2__B1	0.49	0.27	-0.13	0.53	0.88	1.80

Visualize Individual Differences


We can create plots of the trajectory of affect implied by the estimated model parameters, conditional on ± 1 standard deviation and median of any time-independent predictors included in the model using the `ctPredictTIP` function.

The first set of plots created, stored under the `$Process` subobject, includes both latent process and predicted observation trajectories of partner 1 (Affect1) and partner 2 (Affect2). We can see that couples that spent more time together have a faster rate of growth in affect and also plateau at a higher affect value. We plot the latent trajectories only here:

```
tipredplots <- ctPredictTIP(f_indiff, plot=T, tipreds = c('TogetherZ')) #create list of plots
print(tipredplots$Process$Latent) #latent process predictions
```

Figure 19

Predicted Latent Affect Trajectories by Time Spent Together (TogetherZ).




The model implied dynamics, and how these differ with regard to the covariate, are stored under the `$Dynamics` subobject. The first of these shows impulse response functions for an independent shock to each process, and the second shows a shock that is correlated according to the estimated system noise correlation. Here we can visualize the fact that when a person experiences a shock to their affect (increase of 1) their partner is expected to change more over time if they spend more time together, on average — we might interpret this as reflecting their lives being more intertwined so they influence each other more. We visualise the independent shock type here, reflecting what our estimated model suggests should happen under an intervention to one or the other affect process.

```
print(tipredplots$Dynamics$Independent$TogetherZ)
```

Figure 20

Impulse Response Functions for Independent Shocks by Time Spent Together.



Couple Dynamics with Individual Differences and Time-Dependent Moderation (Continuous Time)

We may feel uncomfortable relying on people's retrospective estimate of how much time they spend together to moderate their degree of interdependence. However, we also do not want to assume that couples continuously affect each other even when they are not in each other's presence. For instance, according to our model, if Alex gets fired, Alex's change in affect would immediately influence Robin's affect, even if Robin is completely unaware of this event. If we have information about Alex and Robin's proximity, we can use this information to tell our model to only allow Alex's affect to impact Robin's rate of change when the couple is together. This can be done by moderating our couple's cross-effect parameter by a time-dependent moderator of presence.

Beyond couple dynamics, the same logic applies whenever an observed time-varying condition may turn a dynamic relation on or off, or strengthen and weaken it, from one occasion

to the next. A few examples are: stress exposure might allow symptoms to influence each other only during high-stress periods; medication status might change how quickly a process returns to baseline; poor sleep might increase inertia in mood from one day to the next; social contact might determine whether one person's state can affect another's rate of change; task condition might reveal cross-process coupling only in certain experimental phases; and environmental demands might alter growth or regulation parameters as workload shifts. In each case, the moderator is observed at the same time scale as the process, and it enters the model by scaling a dynamic parameter rather than by adding a direct impulse to the state (as with the therapy example earlier). This is not limited to cross-effects: time-dependent moderators can in principle act on auto-effects, intercepts, cross-effects, or other system parameters. A different extension, not developed in detail here but possible with *ctsem* using similar syntax, would be state-dependent moderation, where the strength of one process's effect depends on the current level of another modeled process – for example, negative affect might influence anxiety only when arousal is already high.

Simulation

To generate data for a model with a time-dependent moderator of cross-effect state dependence, we first need to generate a time-dependent moderator that indicates when a couple is communicating (0 no communication, 1 = communication). To do this, we will sample a value of 0 or 1 from a binomial distribution (`Comm <- rbinom(n = NSubjects, size = 1, prob = .3)`) to reflect whether or not the couple was in contact over the preceding time period.

Second, we add the time-dependent moderator in our model of affect dynamics. This makes the rate of change for partner 1: `dAffect1 <- A*Affect1 + Across[subi] * Affect2 * Comm + B1[subi]`, where *Across* is multiplied by the time-dependent moderator *Comm* (which is either 0 or 1) at each time step. This ensures there is no cross-effect when the couple has not been communicating.

```
NSubjects <- 20
times <- seq(from=0, to=40, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect1 <- rnorm(n = NSubjects, mean = 5, sd = 2)
t0Affect2 <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.4 #continuous time state dependence
Across <- rnorm(n= NSubjects, mean = .1, sd=.05) #cross-effect

# generate continuous intercepts for each individual,
# assuming high affect individuals may partner with high affect
Bcommon <- rnorm(n = NSubjects, mean = 1, sd = .2)
B1 <- Bcommon + rnorm(n = NSubjects, mean = 2, sd = .3) #partner 1
B2 <- Bcommon + rnorm(n = NSubjects, mean = 2, sd = .3) #partner 2

# DIFFUSIONcov: covariance matrix for system noise
DIFFUSIONcov <- matrix(c(.1, .05, .05, .1), nrow=2, ncol=2)

#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA, NSubjects*Nobs),
  Time = rep(NA, NSubjects*Nobs),
  Affect1 = rep(NA, NSubjects*Nobs),
```


```
Affect2 = rep(NA, NSubjects*Nobs))

Nsteps <- 100 #steps between each observation
row <- 0 #initialize row counter
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1 #increment row counter
    Comm <- rbinom(n=1, size=1, prob=.3) #communication dummy
    if(obsi==1){ #set initial affect
      Affect1 <- t0Affect1[subi]
      Affect2 <- t0Affect2[subi]
    }
    if(obsi>1){ #update state via small steps in time
      for(stepi in 1:Nsteps){ #for each step
        #compute deterministic slopes
        dAffect1 <- A*Affect1+Across[subi]*Affect2*Comm + B1[subi]
        dAffect2 <- A*Affect2+Across[subi]*Affect1*Comm + B2[subi]
        # generate correlated noise scaled by stepsize
        systemNoise <- MASS::mvrnorm(n=1, mu=c(0,0),
          Sigma=DIFFUSIONcov/Nsteps)
        #update states using slope and time step, add noise
        Affect1 <- Affect1 + dAffect1/Nsteps + systemNoise[1]
        Affect2 <- Affect2 + dAffect2/Nsteps + systemNoise[2]
      }
    }
    data$Affect1[row] <- Affect1 #input affect data
    data$Affect2[row] <- Affect2 #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Comm[row] <- Comm #input communication dummy
    data$Subject[row] <- subi #input subject data
  }
}

# add measurement error
data$Affect1 <- data$Affect1 + rnorm(n=nrow(data), mean=0, sd = .05)
data$Affect2 <- data$Affect2 + rnorm(n=nrow(data), mean=0, sd = .05)
#plot trajectories with communication moderator
plot_trajectory(data[data$Subject==1,], y="Affect1", y2="Affect2")+
  geom_segment(aes(x=Time, xend=Time, y=0, yend=Comm, linewidth=.5),
    color="red")
```

Figure 21

Affect Trajectories for Both Partners in One Couple with Time-Dependent Communication Moderation.



Model Fitting

To specify a model with a time-dependent moderator in *ctsem* we deviate a little from the earlier example where the predictor had a direct impact on the process itself — here the predictor is only impacting the parameters of the system (the cross effect). First, we specify the name of our time-dependent predictor variable using *TDpredNames*

= c("Comm"). Second, we define the main effect of our predictor using the `TDPREDEFFECT` argument. By default this is a freely estimated parameter. However, in our case we only want the moderation effect, so we set this to 0 (`TDPREDEFFECT = 0`). Third, we modify our drift matrix. Here, we are not going to exactly match our simulation, where we know for certain there is no cross-effect when the couple has a zero on the `Comm` predictor — we will also estimate a baseline parameter, partly serving as a verification of our simulation, and partly as a demonstration of possibilities. So in the `DRIFT` off diagonal elements we specify that our cross-effect is a combination of a baseline parameter (i.e., the cross effect when our moderator is 0) and an interaction term where the `Comm` predictor is directly referenced and multiplied by a parameter that represents the change in cross-effect when `Comm` is 1 — when the couple has been together for the past time period. So our cross-effect becomes `cross12base + cross12mod*Comm`. Sometimes when building more complex structures, it is important to think about where simplifications can be made — in this case, we reduce the number of parameters to be estimated by setting the `auto` and cross-effect parameters of each partner to be equal to each other, by using the same parameter names for both. This illustrates a general modelling tradeoff: additional flexibility can represent richer theory, but equality constraints or other simplifications can be necessary or useful for interpretability and identification.

`ctsem` offers a lot of flexibility in terms of equations nested within each system matrix element. To make use of this flexibility, we need to explicitly tell `ctsem` which text strings represent parameters to be estimated in any complex equations — other text in the string is then treated as mathematical expressions (e.g., `log`, `exp`, `+`, `-`, `*`, `/`, etc.). For this we use the `PARS` argument to tell `ctsem` which text strings are parameters to be estimated, and in this case, also that they are randomly varying across subjects. This gives us entries for `PARS` such as `cross12base||TRUE`. Note that for the ‘simple’ kinds of parameters we have used up until now, we do not need to specify `PARS` at all — only when the parameters are nested within more complex equations do we need to specify `PARS` arguments.

```
m_discmod <- ctModel(
  manifestNames = c("Affect1", "Affect2"), #observed variables
  latentNames = c("Affect1", "Affect2"), #names of latent processes
  TDPredNames = c("Comm"), #time dependent predictors
  time = "Time", #name of time column in dataset
  id = "Subject", #name of subject column in dataset
  type="ct", #use continuous time
  MANIFESTVAR = c('residualSD1', 0,
    0, 'residualSD2'), #sd's of the residual / measurement error
  LAMBDA = diag(1,2), #factor loading matrix
  MANIFESTMEANS=c(0,0), #no measurement intercept / offset
  CINT=c('B1||TRUE', 'B2||TRUE'), #cint with random effects
  TOMEANS=c('t0Affect1||TRUE', 't0Affect2||TRUE'), #initial affect
  TDPREDEFFECT = 0, #disable main effect of comm predictor
  DRIFT = c( #cross effects now moderated and equal to each other
    'auto1', 'cross12base + cross12mod*Comm',
    'cross12base + cross12mod*Comm', 'auto1'),
  DIFFUSION = c( #system noise matrix
    'systemNoise1', 0, #sd for subj 1, 0 in upper triangle
```

```
'systemNoiseCor', 'systemNoise2'),#correlation, sd for subj 2
  PARS=c('cross12base||TRUE', 'cross12mod||TRUE'))#extra params
```

```
f_discmod <- ctStanFit(datalong = data, ctstanmodel = m_discmod)
s_discmod <- summary(f_discmod)
print(s_discmod$popmeans)
```

	mean	sd	2.5%	50%	97.5%
t0Affect1	4.8345	0.3699	4.1571	4.8283	5.5582
t0Affect2	4.6513	0.2702	4.1152	4.6636	5.1685
auto1	-0.3986	0.0175	-0.4336	-0.3977	-0.3646
systemNoise1	0.3166	0.0140	0.2902	0.3160	0.3438
systemNoiseCor	0.1430	0.0123	0.1177	0.1427	0.1677
systemNoise2	10.6870	0.1950	10.2912	10.6895	11.0505
residualSD1	0.0705	0.0277	0.0314	0.0660	0.1336
residualSD2	0.0798	0.1275	0.0044	0.0398	0.4825
B1	3.0556	0.1429	2.7852	3.0523	3.3398
B2	2.8762	0.1381	2.6068	2.8744	3.1448
cross12base	0.0014	0.0178	-0.0362	0.0015	0.0339
cross12mod	0.1051	0.0095	0.0874	0.1049	0.1241

Now we fit and summarise the results (in this case, the model may need to be compiled before fitting, which occurs automatically and usually takes a couple of minutes). The summary shows us that the `cross12base` parameter is estimated to be approximately 0, and the `cross12mod` parameter is estimated to be approximately 0.1300874, 0.0051305, 0.1043938, 0.1157113, 0.0656906, 0.072481, 0.1000715, 0.1345032, 0.1089464, 0.1164402, 0.0235235, 0.1740402, 0.0647841, 0.0805659, 0.1175305, 0.1241835, 0.1881049, 0.115656, 0.0416944, 0.1248069, as we specified in the simulation. This means that when the couple has been communicating for the past time period, the cross-effect is estimated to be approximately 0.15, there is influence from one to the other, otherwise there is no influence.

Different Types of Inputs — Nonlinear Change with a Persistent External Shock (Continuous Time)

The final approach we will consider in terms of model structure is how we can add additional processes to the system that are not directly observed, but can be used to structure the system dynamics. For this example we will turn back to the example where we used therapy as an input to the system, but rather than assume therapy impacts the system and the effect transmits through time with the same dynamics as the rest of the system, we will allow for the persistence, or ‘shape’ of the effect to differ from the rest of the system. This approach can be generalized to achieve things such as separated trends and dynamics processes (Driver & Tomasik, 2023), cross-effects that operate at different time scales (Haehner et al., 2025), oscillatory or more complex dynamic patterns (Voelkle & Oud, 2013), and more.

Our original model implies that the therapy effect is immediate and persists according to the system dynamics. This is probably the simplest approach to thinking about such effects, but in reality, the effects of an intervention or event may be more complex, and may depend on the individual, the context, may induce other changes in the system, and may persist or transfer through the system in different ways to the

typical system behavior. For example, while the regular ups and downs of life may dissipate relatively quickly, a traumatic event may have direct effects that last over a long period and may require more consideration.

To demonstrate this approach, instead of relying on the more ‘general’ system parameters to capture the long term trajectory in response to couples therapy sessions — remembering that from the beginning of this tutorial we have always worked with a rising affect over time — we can now explicitly model this rise, and maintenance at a new level, as a function of therapy sessions. For more examples of intervention effects and how to specify them in *ctsem* we refer the reader to (Driver & Voelkle, 2018b). A major benefit of this approach is that, for example, patients may have very different intervention schedules — allowing the trend to develop based on the specific schedule, rather than a general assumption of a specific form, can be more realistic and flexible. Similar persistent-input structures can represent other psychological processes such as cumulative practice, treatment dose, habit formation, sustained exposure, or gradual learning.

To implement this approach, we need to specify a latent process which is impacted by our observed intervention variable, and is *not measured by any variables*. This process is used to represent the hypothetical development of the intervention effect over time. To try to make this explicit:

Simple impulse form: Intervention directly impacts affect process, dissipation governed by system dynamics.

Separated form: Intervention affects a separate latent process that is independent of the main system dynamics, with either estimated or fixed dynamic parameters — this process in turn then impacts the affect process.

This latent intervention process can have its own autoregressive parameter which allows us to specify an accelerating, dissipating, or stable effect across time. That is, if the state dependence is set to 0, our latent intervention process will remain at a stable value (e.g., remaining at 1 after the initial upwards shock) and thereby generate a permanent shift over time. If the state dependence is set to a negative value (e.g. -0.1), then the effect of the intervention will dissipate over time, but not as quickly as it would have in the pure impulse form.

Simulation

In our data generating block, we include an input variable *M* which directly influences the state of our latent intervention process *TherapyAcc* (accumulated therapy), which is included in much the same way as additional latent processes were previously simulated. The rate of change in our latent intervention process *dTherapyAcc* is given by the previous value of *TherapyAcc* multiplied by a state dependence coefficient, $dTherapyAcc \leftarrow AM * TherapyAcc$. Because the state dependence coefficient is 0 (autoregression = 1) in the case of a persistent input effect, the rate of change is always

0. For simplicity, at the time of therapy sessions the input effect *M* is added directly to *TherapyAcc* — there is no need to track this via the rate of change, as it is an instantaneous impulse.

Between therapy sessions, because the state dependence parameter is zero, the state of the latent intervention process *TherapyAcc* will not change; the therapy sessions generate a step-like process that reflects a persistent input process, gradually accumulating over time as therapy sessions occur (one could argue that it should probably also slowly decay over time, but we will leave that for now). We nevertheless show the code for updating the latent intervention process with a small step in time, to illustrate the process.

To model the impact of our latent intervention process on our affect process, we make the rate of change in affect dependent on the state of the intervention process at the current moment ($dAffect \leftarrow A * Affect + B + TherapyAcc$) — this is just another cross-effect in our drift matrix, from the matrix equation perspective.

```
NSubjects <- 20
times <- seq(from=0, to=100, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.1 #continuous time state dependence
B <- 0.5 #continuous intercept reduced, growth via therapy effect
G <- .1 #system noise coefficient
M <- 0.2 #input variable
AM <- 0 #latent input process autoregression

# randomly sample 3 time points for therapy for each subject
therapyTimes <- lapply(1:NSubjects, function(x){
  return(sample(times, size = 3, replace = FALSE))
})

#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA,NSubjects*Nobs),
  Time = rep(NA,NSubjects*Nobs),
  Affect = rep(NA,NSubjects*Nobs))

Nsteps <- 100 # steps between each observation
row <- 0 #init current row tracker
for(subi in 1:NSubjects){#for each subject
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1 # increment row counter
    #create therapy dummy:
    Therapy <- as.integer(times[obsi] %in% therapyTimes[[subi]])
    if(obsi==1){ #if first time point
      Affect <- t0Affect[subi] #set initial affect
      TherapyAcc <- 0 # initialize latent input process
    }
    if(obsi>1){ #otherwise update state via small steps in time
      for(stepi in 1:Nsteps){ #for each step
        # compute deterministic slopes
        dTherapyAcc <- AM*TherapyAcc # therapy slope update
        dAffect <- A*Affect+B+TherapyAcc #affect slope update
        #update states using slope and time step, add noise
        TherapyAcc <- TherapyAcc + dTherapyAcc/Nsteps
        Affect <- Affect + dAffect/Nsteps +
          G * rnorm(n=1, mean=0, sd=sqrt(1/Nsteps)) #system noise
      }
    }
    TherapyAcc <- TherapyAcc + M*Therapy # effect of therapy
    #store data:
    data$Affect[row] <- Affect #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Subject[row] <- subi #input subject data
    data$Therapy[row] <- Therapy #input therapy dummy
    data$TherapyAcc[row] <- TherapyAcc #input intervention process
  }
}
```


```

}
#add measurement error
data$Affect<-data$Affect+rnorm(n=nrow(data),mean=0,sd=.05)
# plot therapy accumulation process and therapy dummies
plot_trajectory(data[data$Subject %in% 1:4, ],
  color_var = "as.factor(Subject)", y='TherapyAcc') +
  geom_segment(aes(x=Time,xend=Time,y=0,yend=Therapy,linewidth=.5))

```

Figure 22

Latent Intervention Process (TherapyAcc) for Four Subjects with Therapy Sessions.



The plot shows the latent intervention process *TherapyAcc* for the first four subjects in our dataset. We can see that at each therapy session (indicated by the vertical lines), the latent intervention process jumps up by 0.2 (the value of *M*), and then remains stable until the next therapy session.

Model & Fitting

We want to specify a latent intervention process that only reflects the time course of our observed intervention variable, without any extra noise or baseline effects. To do this, we set all the extra influences on the latent intervention process, that do not come directly from the observed intervention variable, to 0. Thus, we set the random fluctuations (*DIFFUSION*) to zero for our latent intervention process. We also set its starting value to 0, by fixing the *TOMEANS* and *TOVAR* to zero for our latent intervention process. We have not had to deal with the *TOVAR* argument before — it is only relevant when there are no individual differences in the starting value of the process, which occurs either when the differences are disabled, or when the starting value is fixed, as in this case. The *TOVAR* matrix provides variance at the initial time point for each subject. We do not want this variance, so we can specify a single zero (which will actually create an appropriate size matrix of zeros). Finally, we fix the therapy processes continuous intercept *CINT* to zero since there is no constant growth in the latent intervention process.

Next, we set the cross-effect of our latent intervention process to 1 using the *DRIFT* matrix (adding a 1 in row 1, column

2 of the matrix). This, specifies a direct effect of the current state of the latent intervention process on the rate of change in our affect process.

Lastly, we set the effect of the observed intervention variable on our latent affect process to zero, and estimate its direct effect on our latent intervention process (*TDPREDEFFECT* = *c(0, "TherapyEffect")*). Thus, our observed intervention directly impact the latent intervention process, which in turn directly impacts our latent affect process. Instead of estimating the effect of *Therapy* on *TherapyAcc*, alternatively we could have fixed this to 1 and estimated the effect of *TherapyAcc* on *Affect*.

```

m_persist <- ctModel(
  manifestNames = "Affect", #observed variables
  latentNames = c("Affect","TherapyAcc"), #latent process names
  TDPredNames = 'Therapy', #time dependent predictors
  time = 'Time', #time column
  id = 'Subject', #subject column
  type='ct', #continuous time
  MANIFESTVAR = 'residualSD', #residual variance
  LAMBDA = matrix(1,nrow=1,ncol=2), #factor loading matrix
  MANIFESTMEANS=0, #no measurement intercept
  CINT=c('B||FALSE', #cint with *no* random effects for Affect,
    "0||FALSE"), # and no CINT at all for therapy accumulation
  TDPREDEFFECT = c(0,"TherapyEffect"), #Therapy on TherapyAcc
  TOMEANS=c("t0Affect||TRUE", #initial affect with random effects,
    "0||FALSE"), # and therapy accumulation starts at 0
  TOVAR=0, # no within-subject variance in starting values
  DRIFT = c( #auto effect and TherapyAcc effect on affect.
    'stateDependence', 1, #autoeffect
    0, 0), #therapy accumulation process is static
  DIFFUSION = c('systemNoise',0, #system noise for affect
    0, 0)) #no system noise for therapy accumulation

```


```
f_persist <- ctStanFit(datalong = data, ctstanmodel = m_persist)
```

We can plot the smoothed estimates of the latent variable states for the first three subjects in our dataset, which shows the increase in the accumulated therapy at the time of the intervention, which causes increases in affect over time. Here we can see the instantaneous impulse of the therapy session on accumulated therapy, and the long term growth in affect due to accumulated therapy.

```
ctKalman(f_persist, plot=TRUE, subjects=1:3, kalmanvec='etasmooth')
```

Figure 23

Smoothed Latent Processes for Three Subjects with Accumulating Therapy Input.



Model Evaluation and Comparison

To end this tutorial, we are going to investigate how to evaluate and compare different model structures when we do not know the structure of our observed data. In empirical settings, this can be a very difficult and fascinating problem that requires deep inquiry, and we cannot cover all aspects and nuance here. To get people started with model evaluation, we will demonstrate some tools and approaches that can help determine whether a model provides a good representation of the data, and whether it is better than any competing explanations (models) we may have.

There is no automatic rule for choosing a dynamic model. Model choice should triangulate theory, the relevant time scale, data density, identifiability, model-implied behavior, diagnostics, and parsimony. A more flexible model will often improve within-sample summaries, but this does not mean the extra parameters are interpretable or supported. Conversely, a model that is too simple may leave distinct sources of variation bundled into one parameter. The examples below therefore use R^2 -style summaries only as one descriptive starting point, and then focus on prediction plots, residual autocorrelation, and posterior-predictive-style checks.

So let's imagine that after all the model building we have done, we still think that the best representation of an affect process is a linear trajectory. If we have data from 20 patients over 21 weeks, we can use the R^2 estimate to check how much variance in the data has been explained by the model. For our example we will generate data from a known model (simplifying the explanations) but normally the true model is not known!

Generating Data

To generate our data we will recycle our code block from earlier (see 'Simulation: Nonlinear Change with Random Fluctuations (Continuous Time)').

```
NSubjects <- 20
times <- seq(from=0, to=20, by=1) #time points of measurement
Nobs <- length(times) #number of observations per subject
t0Affect <- rnorm(n = NSubjects, mean = 5, sd = 2)
A <- -.1 #continuous time state dependence
B <- 1 #continuous intercept
G <- 0.2 #system noise coefficient

#create empty data.frame to fill step by step
data <- data.frame(Subject= rep(NA,NSubjects*Nobs),
  Time = rep(NA,NSubjects*Nobs),
  Affect = rep(NA,NSubjects*Nobs))

Nsteps <- 100 #steps between each observation
row <- 0 #init current row tracker
for(subi in 1:NSubjects){
  for(obsi in 1:Nobs){ #for each observation of a subject
    row <- row + 1
    #if first time point, set to initial affect
    if(obsi==1) Affect <- t0Affect[subi]
    if(obsi>1){ #else compute new state via small steps in time
      for(stepi in 1:Nsteps){ #every step, do...
        dAffect <- A*Affect + B #compute affect slope
        Affect <- Affect + #new state = old state +
          dAffect * 1/Nsteps + #deterministic change +
          G * rnorm(n=1, mean=0, sd=sqrt(1/Nsteps)) #random change
      }
    }
    data$Affect[row] <- Affect #input affect data
    data$Time[row] <- times[obsi] #input time data
    data$Subject[row] <- subi #input subject data
  }
}
#add measurement error
data$Affect <- data$Affect+rnorm(n=nrow(data), mean = 0, sd = .05)
```

Fitting the Wrong Model: Linear Growth

We now specify and fit our linear growth model, with no state-dependency and no system noise, in ctsem.

```
m_linear <- ctModel( #define the ctsem model
  manifestNames = "Affect", #names of observed variables in dataset
  latentNames = "Affect", #names of latent processes
  time = 'Time', #name of time column in dataset
  id = 'Subject', #name of subject column in dataset
  type='ct', #continuous time
  MANIFESTVAR = 'residualSD', #sd of residual / measurement error
  LAMBDA = matrix(1,nrow=1,ncol=1), #factor loading matrix
  MANIFESTMEANS=0, #no measurement intercept
  CINT='B||FALSE', #continuous intercept with *no* random effects
  TOMEANS='t0Affect||TRUE', #initial affect with random effects
  DRIFT = 0,
  DIFFUSION = 0)
```

```
f_linear <- ctStanFit(dataalong = data, ctstanmodel = m_linear)
```

A common form of model evaluation is to consider how much of the variance in the data has been explained by the model. This is typically quantified using the R^2 statistic (Nakagawa & Schielzeth, 2013) or some variation thereof. The R^2 tells us how much of the total variance in our data is explained by our model by dividing the residual (unexplained) variance ($V\hat{y}$) by the total variance (Vy) and subtracting this from 1. Because in ctsem models the total variance is not necessarily constant across individuals or time, this is computed using standardized residuals, and reported (as standardized residual variance / covariance) at the top of the ctsem summary output. This should not be treated as a stand-alone model selec-

tion criterion: complex models can improve within-sample explained variance while still implying the wrong process.

```
s=summary(f_linear) #store summary of the fit
print(s$residCovStd) #print standardized residual (co)variance
```

```
Affect
Affect 0.236
```


Now we can use the standardized residual variance estimate (reported at the top of summary under ‘residCovStd’, equivalent to $1 - R^2$) to see how much of the variance in our empirical data our model has accounted for. We can see that our model explains a high proportion of the variance in our observations. This is fairly typical for repeated measures data, and doesn’t imply that our model is appropriate, but provides a useful baseline for comparison.

Let’s begin considering model appropriateness by looking at the models forward predictions (i.e., conditional on parameter estimates and earlier data, represented by lines in the plot) against the observed data (represented by dots) for a few subjects:

```
ctKalman(f_linear, plot=TRUE, subjects=1:3,
kalmanvec = c('y', 'yprior')) #y=observed, yprior=predictions
```

Figure 24

Predictions from Misspecified Linear Model Versus Observed Data for Three Subjects.



A first question that might arise here, is why aren’t the model predictions the simple linear trajectory we specified? This occurs because as the model steps through time and ‘sees’ more data, it learns the individual difference parameters for each subject, and thus the predictions adjust and do not follow the simple linear trajectory. If we plotted the ‘smoothed’ predictions (conditional on all data, not just earlier data, use `ysmooth` instead of `yprior`) they would follow the linear trajectory. The forward predictions we see here, conditional only on parameter estimates and past data, are what the likelihood calculations and residuals are based on.

If using an appropriate model, the residuals here — the difference between the observed and predicted data — should be entirely *unpredictable* based on earlier data or residuals.


While far from a formal test, a visual inspection of the model predictions versus observed data suggests that this is *not* the case, as we can see some clear structure in the observed data that is not being captured by the model. This is visible in that when the model prediction is too low at one observation, it is also likely to be too low at the next observation — our residuals are *autocorrelated*. Autocorrelated residuals indicate that our model leaves some predictable structure unexplained (Gelman et al., 1996) and we could modify our model to improve the predictions. This state of affairs is broadly referred to as *model misspecification*. The predictions may still be useful for some purposes, but we should be extremely careful about interpretations of the model parameters or other generalisations.

An approach to examine all residuals at once in a more comprehensive manner is to check their autocorrelation. We can create a plot of the autocorrelation between residuals to visualize whether our model is missing any systematic patterns in our data. To do this we use the `ctACFresiduals` function.

```
ctACFresiduals(f_linear) #plot autocorrelation of residuals
```

Figure 25

Autocorrelation of Residuals from the Misspecified Linear Model.




From the plot we can see that there is a high autocorrelation between residuals, about 0.6 between temporally adjacent residuals. So now we know two things: (1) Our model explains a substantial amount of the variance in our data, (2) there is still room for improvement. The natural follow-up question is what should we improve? To get an idea of model misspecification we can generate data based on the parameter estimates of our model and plot those against our observed data. This approach is analogous to posterior predictive checks in a Bayesian framework, but can be applied in

non-Bayesian contexts as well (Gelman & Shalizi, 2013).

```
f_linear <- ctStanGenerateFromFit(f_linear, nsamples=200,
  fullposterior = F, cores = 2) #add generated data to fit object
postpred<-ctPostPredPlots(f_linear) #create plots
print(postpred[[1]]) #show first plot --- overall density
```

Figure 26

Posterior Predictive Check: Density of Likelihood and Variables (Linear Misspecification).




The density-based posterior predictive check compares the empirical distribution of the observed variables and likelihood contributions with distributions generated from the fitted model. It is difficult to get a clear sense of exactly where things are imperfect from such a plot, but this aggregate view provides another approach to see that there clearly are some patterns in the data that deviate from the model's expectations. Perfect overlap between the empirical and model-implied distributions for each variable doesn't necessarily mean the model is appropriate, but does at least indicate that the overall distributions are similar.

```
print(postpred[[2]]) #show second plot --- expectations over time
```

Figure 27

Posterior Predictive Check: Model-Implied Versus Empirical Expectations over Time (Linear Misspec).



The trajectory-based posterior predictive check shows the model implied expectations (red) versus the empirical data (black) over time. Here it is quite clear that we are fitting a linear model to a nonlinear trajectory. Our model also seems to underestimate the variance at earlier time points while overestimating the variance at later observations. But how can we improve our model to better account for these observations?

Taking a closer look at the observed individual growth trajectories (fainter black lines), we can see there seems to be a dependency between people's initial affect values and their rate of change in affect. That is, people who start with very high affect values have an almost flat growth curve, while people with a low initial level have a very steep growth curve which slowly plateaus as they reach higher values of affect. Thus, it seems that we should add a state-dependency parameter to our model to capture this relationship. In empirical applications, this step should be justified by theory and diagnostics together: state dependence is most defensible when it is plausible that the current state at least partially determines subsequent change.

Fitting a Better Model: Adding State-Dependence in Continuous Time

```
m_statedep <- ctModel(
  manifestNames = "Affect", #observed variables
  latentNames = "Affect", #latent processes
  time = 'Time', #time column
  id = 'Subject', #subject column
  type='ct', #continuous time
  MANIFESTVAR = 'residualSD', #residual variance
  LAMBDA = matrix(1,nrow=1,ncol=1), #factor loading matrix
  MANIFESTMEANS=0, #no measurement intercept
  CINT='B||FALSE', #continuous intercept with *no* random effects
  TOMEANS='t0Affect||TRUE', #initial affect with random effects
  DRIFT = 'stateDependence', #state-dependence
  DIFFUSION = 0) #no system noise
```

```
f_statedep <- ctStanFit(datalong = data, ctstanmodel = m_statedep)
```

```
s=summary(f_statedep) #store summary of the fit
print(s$residCovStd) #print standardized residual (co)variance
```


```
Affect
Affect 0.101
```

Now we can again see if adding this state-dependence term has helped us explain more variance in our data by considering the R^2 — and indeed, we can see that we have around 10% residual variance, or 90% explained variation in our data — noticeably more than before. This is encouraging, but it is not sufficient: we also need to check whether the model-implied behavior and residual structure are more appropriate. So now we can go view our autocorrelation plot to check if there is still information left in the residuals that our model could explain.

```
ctACFresiduals(f_statedep) #plot autocorrelation of residuals
```

Figure 28

Autocorrelation of Residuals from the State-Dependence Model.




We can see that the autocorrelation has also dropped substantially, but there is still a moderate amount of autocorrelation in our residuals. Which is rightly telling us that there is more predictable variation that our model can leverage. So let's try to find out what it is by plotting the model generated data against our observed data.

```
fit <- ctStanGenerateFromFit(f_statedep, nsamples=200,
  fullposterior = F, cores = 2) #add generated data to fit object
postpred <- ctPostPredPlots(f_statedep) #create plots
print(postpred[[1]]) #show first plot --- overall density
```

Figure 29

Posterior Predictive Check: Density of Likelihood and Variables (State-Dependence Model).




Looking at the first plot, the model implied distributions are now quite similar to the empirical distributions, but there is still some difference — it's not so easy to determine if such a difference is important, from this visual.

```
print(postpred[[2]]) #show second plot --- expectations over time
```

Figure 30

Posterior Predictive Check: Model-Implied Versus Empirical Expectations over Time (State-Dependence Model).




The trajectory plot shows that our model implied trajectory almost perfectly fits our observed trajectory, and the variance now also reduces over time in line with the observed trajectories converging to a similar level. So, we know there appears to be some predictable variation that our model could better account for, but it's difficult to determine how to do this. Examining empirical versus model implied lagged correlations can provide more insight:

```
ctFitCovCheck(f_statedep, plot = TRUE, cor=TRUE)
```

\$Affect

Figure 31

Empirical Versus Model-Implied Lagged Correlations (State-Dependence Model).



This plot shows the empirical versus model implied lagged correlations for each variable (in a multivariate scenario, we would also see cross-correlations). We can see that the model implied correlations between one time point and later time points are stable, while the empirical correlations decrease over time. This makes sense because although we've now

specified a state dependence parameter in the model, causing growth to slow and stabilise over time, we didn't allow for any *new* sources of variation to be introduced into the process — the same variation that was there in the beginning persists on and on, smaller in magnitude, but equally correlated with earlier time points. Introducing a system noise term to our model will allow for new sources of variation to be introduced into the process, and we can see how this changes the lagged correlations. Substantively, this step is warranted when the theory and diagnostics suggest that unmodelled influences affect the latent process itself.

```
m_sysnoise <- ctModel(
  manifestNames = "Affect", #observed variables
  latentNames = "Affect", #latent processes
  time = 'Time', #time column
  id = 'Subject', #subject column
  type='ct', #continuous time
  MANIFESTVAR = 'residualSD', #residual variance
  LAMBDA = matrix(1,nrow=1,ncol=1), #factor loading matrix
  MANIFESTMEANS=0, #no measurement intercept
  CINT='B||FALSE', #continuous intercept with *no* random effects
  TOMEANS='t0Affect||TRUE', #initial affect with random effects
  DRIFT = 'stateDependence', #state-dependence
  DIFFUSION = 'systemNoise') #system noise

f_sysnoise <- ctStanFit(datalong = data, ctstanmodel = m_sysnoise)


s=summary(f_sysnoise) #store summary of the fit
print(s$residCovStd) #print standardized residual (co)variance

      Affect
Affect 0.082

ctACFresiduals(f_sysnoise) #plot autocorrelation of residuals
```

Figure 32

Autocorrelation of Residuals from the Model with System Noise.



Overall, we can see that fitting the model to the data it generated leads to a greater R^2 , a near-zero autocorrelation between residuals, and model-implied lagged correlations that map well onto the observed lagged correlations.

In empirical contexts, this process is going to be more complicated — individual differences will likely need more

serious consideration, and very often it will be difficult to impossible to obtain a model that appears perfect while also retaining some level of interpretability. Sensitivity checks, wherein core inferences are checked across a range of model specifications, can help to assess the robustness of the inferences to different model specifications (for an example of the approach in the context of ctsem see [Driver, 2025](#)).

Practical Troubleshooting

Dynamic state-space models can fail in ways that are informative. Common warning signs include Hessian warnings after fitting or extremely large or small standard errors / uncertainty intervals, unstable estimates across small changes in specification, parameter estimates at boundary values, implausibly fast or slow dynamics, and poor posterior predictive behavior. These problems can arise from the data, the model, or their mismatch. For example, variables with little or no variance, sparse observations, low reliability, an unsuitable time unit, or too many random effects can all make estimation fragile.

Several practical checks are useful if encountering a difficult fit. First, inspect the data scale and the time scale. At least within ctsem, models are often easier to estimate when time units of roughly 0.1 to 1.0 represent moderate amounts of change; very large or very tiny intervals can lead to numerical problems and or start values that are too extreme in one direction or the other. Centering or scaling observed variables can also help make start values and unconstrained parameter scales more reasonable. Scaling of covariates is particularly important — this is not handled by the software automatically. Time dependent predictors should be zero except when some effect occurs at a specific time. Time independent predictors should in most cases be centered to have mean zero, such that resulting parameter estimates are interpretable with respect to average values on the predictor variable. They should also have roughly variance of one, though this is not always necessary. A second check is to inspect the expanded model specification to ensure that ctsem defaults have not introduced unintended free parameters or random effects. Third, fit simpler intermediate models and add complexity only when the previous model behaves sensibly. Fourth, sensitivity checks are a useful way to assess the robustness of the inferences to different model specifications — if you are not sure on the appropriateness of a model, it is a good idea to check how key inferences change across a range of model specifications.

Finally, treat warnings as prompts for model criticism rather than as purely technical obstacles. A Hessian warning, boundary estimate, or sensitivity to specification changes may indicate that the data do not adequately identify the model, that the time scale is poorly chosen, or that a theoretically distinct influence is being forced into the wrong parameter because of misspecification. The appropriate response may be to rescale, simplify, collect denser or more reliable

data, or revise the substantive model.

Further Works and Resources

Beyond the core `ctsem` workflow, several tutorials and reviews map the broader longitudinal modelling landscape and can help with the array of choices to be made in linking theory to models and data. For discrete-time foundations, a concise, practice-oriented overview of growth curve models addresses common pitfalls and implementation details (Curran et al., 2010). For intensive longitudinal data that blur boundaries between multilevel models, time series, and SEM, dynamic structural equation modelling (DSEM) provides an integrated framework with worked examples and diagnostics (Asparouhov et al., 2018). Bridging discrete- and continuous-time perspectives, step-by-step introductions show how cross-lagged and autoregressive panel models can be translated into their continuous-time counterparts with stochastic differential equations and estimated within SEM (Voelkle et al., 2012). For researchers using cross-lagged panel models specifically, some criticisms clarify interpretational limits and motivates model choices that better align statistical parameters with substantive within-person processes (Berry & Willoughby, 2017; Hamaker et al., 2015).

For dynamic-systems modelling that formalizes change mechanisms directly as differential (or stochastic differential) equations, there are other accessible software and methodological resources. The `dynr` package paper provides a hands-on introduction to specifying linear/nonlinear, discrete/continuous-time models with code recipes and estimation details (Ou et al., 2019). Methodological tutorials demonstrate hierarchical stochastic differential equation models for affect dynamics, highlighting how drift and diffusion parameters map to psychological regulation and noise (Oravecz et al., 2011), and how short and long term timescales may be considered in the one model (Boker et al., 2024).

Linking theory to model requires explicit mappings from theoretical constructs to parameters, along with clearly stated scope conditions and levels of analysis, some general guidelines of which can be found in Wilson and Collins (2019). A range of works explore how formal models can serve, inform, or be tested against theories, and advocates using formal models as a practical toolkit for theory construction (Haslbeck et al., 2022; Robinaugh et al., 2021). Recent proposals to integrate intra- and inter-individual phenomena make these mappings explicit, clarifying which parameters encode within-person dynamics versus between-person structure and how each bears on theoretical claims (Borsboom & Haslbeck, 2024). For panel designs, careful attention to time scale and confounding is crucial: cross-lagged effects are often misinterpreted when temporal assumptions are mismatched or when time-invariant confounders are not handled appropriately; guidance is available for diagnosing these

issues and providing more defensible causal interpretations (Driver, 2025; Rohrer & Murayama, 2023). Finally, when theories posit continuous change, formalizing them as continuous-time systems can yield parameters (e.g., drift, diffusion) that align more directly with hypothesized mechanisms and developmental relations (Driver & Tomasik, 2023).

Conclusion

Modern computational tools and software packages such as `ctsem` simplify the process of building intricate statistical models to meet the surge of interest in intensive longitudinal data. However, the interpretation of complicated models remains challenging despite the increasing ease of their implementation. The central lesson of this tutorial is therefore not simply how to fit a particular software model, but how to move between theory, generative assumptions, simulation, estimation, model checking, and revision.

The running example focused on couples' affect dynamics, but the same modelling logic applies whenever psychological theory concerns change, regulation, coupling, inputs, or heterogeneity over time. State dependence can represent inertia or recovery; cross-effects can represent influence among people, symptoms, behaviors, or physiological processes; diffusion can represent unmodelled inputs to the latent process; measurement error can represent imperfect observation; and moderators can represent stable or time-varying conditions under which dynamics differ. These ideas are useful only to the extent that the model, data, and theory support the interpretation. A complex model that is weakly identified can mislead, but so can an overly simple model that bundles distinct mechanisms into one parameter.

There are many more extensions and model variants that could be considered, but two are worth noting because they follow naturally from the models we have specified and demonstrate how the concepts can be applied to more complex scenarios. First, one could allow state-dependent moderation, where the effect of one process on another depends on the current level of one or both processes. For example, stress might increase symptom coupling only when arousal is already high, or social support might buffer negative affect only below a certain level of distress. Second, one could use nonlinear measurement models when the observed scale changes sensitivity across the latent range, such as we might see if for instance negative affect tends to change only slowly at lower levels and more rapidly at higher levels, leading to skewed distributions of the observed variable but a latent process that can still be modelled as continuous with Gaussian noise.

The practical goal behind this tutorial is to provide the tools and intuition for iterative and cautious model building. Researchers should begin with explicit theoretical assumptions, simulate the implications of those assumptions where possible, fit models that are no more complex than the the-

ory and data can support, check whether the model-implied behavior is plausible, and simplify or revise when the model fails. We hope this tutorial helps researchers use dynamic systems models as transparent tools for theory development rather than as black-box descriptions of longitudinal data.

References

- Aalen, O. O., Røysland, K., Gran, J. M., Kouyos, R., & Lange, T. (2016). Can we believe the DAGs? A comment on the relationship between causal DAGs and mechanisms. *Statistical Methods in Medical Research*, 25(5), 2294–2314. <https://doi.org/10.1177/0962280213520436>
- Asparouhov, T., Hamaker, E. L., & Muthén, B. (2018). Dynamic Structural Equation Models. *Structural Equation Modeling: A Multidisciplinary Journal*, 25(3), 359–388. <https://doi.org/10.1080/10705511.2017.1406803>
- Bates, D., Mächler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1), 1–48. <https://doi.org/10.18637/jss.v067.i01>
- Berry, D., & Willoughby, M. T. (2017). On the practical interpretability of cross-lagged panel models: Rethinking a developmental workhorse. *Child Development*, 88(4), 1186–1206. <https://doi.org/10.1111/cdev.12660>
- Boker, S. M., Daniel, K. E., & Orzek, J. (2024). Separating Long-Term Equilibrium Adaptation from Short-Term Self-Regulation Dynamics Using Latent Differential Equations. *Multivariate Behavioral Research*, 59(6), 1177–1187. <https://doi.org/10.1080/00273171.2023.2228302>
- Borsboom, D., & Haslbeck, J. (2024). Integrating Intra- and Interindividual Phenomena in Psychological Theories. *Multivariate Behavioral Research*, 59(6), 1290–1309. <https://doi.org/10.1080/00273171.2024.2336178>
- Borsboom, D., van der Maas, H. L. J., Dalege, J., Kievit, R. A., & Haig, B. D. (2021). Theory construction methodology: A practical framework for building theories in psychology. *Perspectives on Psychological Science*, 16(4), 756–766. <https://doi.org/10.1177/1745691620969647>
- Butler, E. A., & Randall, A. K. (2013). Emotional coregulation in close relationships. *Emotion Review*, 5(2), 202–210. <https://doi.org/10.1177/1754073912451630>
- Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M., Guo, J., Li, P., & Riddell, A. (2017). Stan: A probabilistic programming language. *Journal of Statistical Software*, 76(1). <https://doi.org/10.18637/jss.v076.i01>
- Curran, P. J., Obeidat, K., & Losardo, D. (2010). Twelve frequently asked questions about growth curve modeling. *Journal of Cognition and Development*, 11(2), 121–136. <https://doi.org/10.1080/15248371003699969>
- Driver, C. C. (2023). Inference with cross-lagged effects – common causes in dynamic systems. *PsyArXiv*. <https://doi.org/10.31234/osf.io/y3e9d>
- Driver, C. C. (2025). Inference with cross-lagged effects—Problems in time. *Psychological Methods*, 30(1), 174–202. <https://doi.org/10.1037/met0000665>
- Driver, C. C., Oud, J. H. L., & Voelkle, M. C. (2017). Continuous Time Structural Equation Modeling with R package ctsem. *Journal of Statistical Software*, 77(5), 1–35. <https://doi.org/10.18637/jss.v077.i05>
- Driver, C. C., & Tomasik, M. J. (2023). Formalizing developmental phenomena as continuous-time systems: Relations between mathematics and language development. *Child Development*, 94(6), 1454–1471. <https://doi.org/10.1111/cdev.13990>
- Driver, C. C., & Voelkle, M. C. (2018a). Hierarchical Bayesian continuous time dynamic modeling. *Psychological Methods*, 23(4), 774–799. <https://doi.org/10.1037/met0000168>
- Driver, C. C., & Voelkle, M. C. (2018b). Understanding the time course of interventions with continuous time dynamic models. In K. van Montfort, J. H. L. Oud, & M. C. Voelkle (Eds.), *Continuous time modeling in the behavioral and related sciences* (pp. 79–109). Springer International Publishing. <http://www.springer.com/de/book/9783319772189>
- Driver, C. C., Voelkle, M. C., & Oud, J. H. L. (2025). *Ctsem: Continuous time structural equation modelling* [Manual]. CRAN. <https://cran.r-project.org/package=ctsem>
- Durbin, J., & Koopman, S. J. (2012). *Time series analysis by state space methods* (2nd ed.). Oxford University Press.
- Epskamp, S. (2020). Psychometric network models from time-series and panel data. *Psychometrika*, 85(1), 206–231. <https://doi.org/10.1007/s11336-020-09697-3>
- Gabry, J., Simpson, D., Vehtari, A., Betancourt, M., & Gelman, A. (2019). Visualization in bayesian workflow. *Journal of the Royal Statistical Society: Series A*, 182(2), 389–402. <https://doi.org/10.1111/rssa.12378>
- Gelman, A., Meng, X.-L., & Stern, H. (1996). Posterior predictive assessment of model fitness via realized discrepancies. *Statistica Sinica*, 6(4), 733–760. <https://www.jstor.org/stable/24306036>
- Gelman, A., & Shalizi, C. R. (2013). Philosophy and the practice of Bayesian statistics. *British Journal of Mathematical and Statistical Psychology*, 66(1), 8–38. <https://doi.org/10.1111/j.2044-8317.2011.02037.x>
- Haehner, P., Driver, C. C., Hopwood, C. J., Luhmann, M., Fliedner, K., & Bleidorn, W. (2025). The dynamics of self-esteem and depressive symptoms across days, months, and years. *Journal of Personality and Social Psychology*. <https://doi.org/10.1037/pspp0000542>
- Hamaker, E. L., Kuiper, R. M., & Grasman, R. P. P. P. (2015). A critique of the cross-lagged panel model. *Psychological Methods*, 20(1), 102–116. <https://doi.org/10.1037/a0038889>
- Harvey, A. C. (1989). *Forecasting, structural time series*

- models and the kalman filter*. Cambridge University Press.
- Haslbeck, J. M. B., Ryan, O., Robinaugh, D. J., Waldorp, L. J., & Borsboom, D. (2022). Modeling psychopathology: From data models to formal theories. *Psychological Methods*, 27(6), 930–957. <https://doi.org/10.1037/met0000303>
- Higham, D. J. (2001). An algorithmic introduction to numerical simulation of stochastic differential equations. *SIAM Review*, 43(3), 525–546. <https://doi.org/10.1137/S0036144500378302>
- Kalman, R. E. (1960). A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1), 35–45. <https://doi.org/10.1115/1.3662552>
- Kenny, D. A. (2005). Cross-Lagged Panel Design. In *Encyclopedia of Statistics in Behavioral Science*. John Wiley & Sons, Ltd. <https://doi.org/10.1002/0470013192.bsa156>
- Kenny, D. A., Kashy, D. A., & Cook, W. L. (2006). *Dyadic data analysis*. Guilford Press.
- Kloeden, P. E., & Platen, E. (1992). *Numerical solution of stochastic differential equations*. Springer.
- Kuiper, R. M., & Ryan, O. (2018). Drawing conclusions from cross-lagged relationships: Re-considering the role of the time-interval. *Structural Equation Modeling: A Multidisciplinary Journal*, 25(5), 809–823. <https://doi.org/10.1080/10705511.2018.1431046>
- Lütkepohl, H. (2005). *New introduction to multiple time series analysis*. Springer.
- Nakagawa, S., & Schielzeth, H. (2013). A general and simple method for obtaining R^2 from generalized linear mixed-effects models. *Methods in Ecology and Evolution*, 4(2), 133–142. <https://doi.org/10.1111/j.2041-210X.2012.00261.x>
- Oravecz, Z., Tuerlinckx, F., & Vandekerckhove, J. (2011). A hierarchical latent stochastic differential equation model for affective dynamics. *Psychological Methods*, 16(4), 468–490. <https://doi.org/10.1037/a0024375>
- Ou, L., Hunter, M. D., & Chow, S.-M. (2019). What's for dynr: A package for linear and nonlinear dynamic modeling in R. *The R Journal*, 11(1), 91–111. <https://doi.org/10.32614/rj-2019-012>
- Robinaugh, D. J., Haslbeck, J. M. B., Ryan, O., Fried, E. I., & Waldorp, L. J. (2021). Invisible hands and fine calipers: A call to use formal theory as a toolkit for theory construction. *Perspectives on Psychological Science*, 16(4), 725–743. <https://doi.org/10.1177/1745691620974697>
- Rohrer, J. M., & Murayama, K. (2023). These are not the effects you are looking for: Causality and the within-/between-persons distinction in longitudinal data analysis. *Advances in Methods and Practices in Psychological Science*, 6(1), 25152459221140842. <https://doi.org/10.1177/25152459221140842>
- Rosseel, Y., & Vidal, M. (2025). A tutorial for understanding SEM using R: Where do all the numbers come from? *British Journal of Mathematical and Statistical Psychology*. <https://doi.org/10.1111/bmsp.70003>
- Ryan, O., & Hamaker, E. L. (2022). Time to intervene: A continuous-time approach to network analysis and centrality. *Psychometrika*, 87(1), 214–252. <https://doi.org/10.1007/s11336-021-09767-0>
- van Geert, P. (2011). The contribution of complex dynamic systems to development. *Child Development Perspectives*, 5(4), 273–278. <https://doi.org/10.1111/j.1750-8606.2011.00197.x>
- Voelkle, M. C., & Oud, J. H. L. (2013). Continuous time modelling with individually varying time intervals for oscillating and non-oscillating processes. *British Journal of Mathematical and Statistical Psychology*, 66(1), 103–126. <https://doi.org/10.1111/j.2044-8317.2012.02043.x>
- Voelkle, M. C., Oud, J. H. L., Davidov, E., & Schmidt, P. (2012). An SEM approach to continuous time modeling of panel data: Relating authoritarianism and anomia. *Psychological Methods*, 17(2), 176–192. <https://doi.org/10.1037/a0027543>
- Wickham, H. (2011). Ggplot2. *Wiley Interdisciplinary Reviews: Computational Statistics*, 3(2), 180–185.
- Wilson, R. C., & Collins, A. G. (2019). Ten simple rules for the computational modeling of behavioral data. *eLife*, 8, e49547. <https://doi.org/10.7554/eLife.49547>